

Colour from Gravity

SU(3) as a Closure Attractor in the Palatini Bracket

The Strong Force Algebra from Vierbein–Connection Geometry and Unitarity

Bradley Wallace

Independent Researcher

Publication status (24 September 2026). This version includes a claim-specific amendment at the end of the main text. That amendment governs earlier physical interpretations and distinguishes retained calculations from withdrawn or unresolved inferences.

Abstract

The colour gauge algebra $\mathfrak{su}(3)$ emerges as a geometric selection effect within the Palatini formulation of gravity. The Asymmetric Codependent System (ACS) framework provides the technical foundation: the net transfer entropy $\Delta\mathcal{I}$ between two mutually constraining fields has the Lie bracket $[f, g]$ as its exact second-order Taylor coefficient (BCH–TE morphism). Applied to the vierbein–connection pair $(e^a{}_\mu, \omega^{ab}{}_\mu)$, the bracket map generates $\mathfrak{sl}(4)$ (rank 15, exact), which the Palatini decomposition splits into a 6-dimensional Lorentz sector and a 9-dimensional torsion sector. The split real form $\mathfrak{sl}(3, \mathbb{R})$ embeds as the only 8-dimensional subalgebra achieving numerically exact closure ($\mathcal{D} < 10^{-14}$) among 50,000 random samples; a chirality map $J(T) = i \text{sym}(T) + \text{anti}(T)$, by Cartan classification, is the unique linear map producing a compact real form. The mechanism $\mathfrak{sl}(3, \mathbb{R}) \xrightarrow{J} \mathfrak{su}(3)$ is verified (all 28 brackets close; all generators skew-Hermitian). The Barbero–Immirzi parameter $\gamma = 0.274$ is recovered from the ACS information-balance condition. All computational claims are backed by the accompanying verification suite. Predictions: a 49 keV sterile neutrino, $\theta_{\text{QCD}} = 0$ exactly, and a torsion coupling hierarchy 0:1:4 across the unbroken, electroweak and colour-lepton generators (on eigenvectors of $\text{ad}_{T_{B-L}}$ this is a two-valued spectrum, 0 and 32/9; see the torsion-hierarchy section).

Contents

1	Introduction	1
1.1	The core observation	1
1.2	Summary of results	1
1.3	Paper structure	1
2	The Asymmetric Codependent System	2
2.1	Measure-theoretic foundation	2
2.2	Definitions	2
2.3	The BCH–transfer-entropy morphism	5
2.4	Exact integer automaton confirmation	7
3	The ACS in Gauge Theory and Gravity	7
3.1	The covariant derivative as the ACS coupling	8
3.2	Gravity as an ACS: the Palatini formulation	8
3.3	Ricci flow as ACS dynamics	9
4	The $\mathfrak{GL}(4)$ Fiber Algebra and the Electroweak Sector	9
4.1	The algebra chain	10
4.2	The image of the asymmetry map	11
4.3	Selection principle: why $\mathfrak{sl}(3, \mathbb{R})$ and no other	12
4.4	Colour charges from Palatini geometry	14
4.5	Electroweak containment	16
4.6	The precise boundary: where $SU(3)$ fails	17
4.7	What remains open	18
5	ACS Applied to Quantum Gravity	19
5.1	Quantum gravity as an ACS problem	19
5.2	The Barbero–Immirzi parameter from ACS information balance	20
5.3	Wheeler–DeWitt as ACS quantum attractor	22
5.4	Black holes and holographic resolution	22
5.5	Predictions	23
6	Open Problems and Phase 2 Roadmap	23
6.1	The mass hierarchy angle (resolved)	23
6.2	The Barbero–Immirzi constrained value	24
6.3	The neutrino mass spectrum (characterised)	24
6.4	Quark Koide and QCD corrections	25
6.5	The Higgs potential from the BCH bracket	25
6.6	Dimensional devolution	27
6.7	Contextualization within current literature	28
A	Proof of Theorem 2.1 (T3)	31
B	Computational Status of Key Predictions	32
C	Status of the CKM Derivation	34
C.1	Torsion coupling hierarchy and vacuum energy cancellation	35
C.2	The irreducible boundary	36
D	Geometric Vision (Plain-Language Summary)	39
E	Experimental Tests	40

1 Introduction

1.1 The core observation

Consider any system in which two objects co-exist without either being prior. The vierbein and the spin connection in the Palatini formulation. The metric and the connection in general relativity. In every case, two structurally different objects are *codependent* — neither can be defined without the other — and *asymmetric* — the kind of information each carries about the other differs in character.

This paper demonstrates that this structural asymmetry, applied to the Palatini formulation of gravity, geometrically selects the colour gauge algebra $\mathfrak{su}(3)$ of the strong nuclear force.

Scope and inputs. The derivation has two stages and two inputs. Stage 1 (algebraic, no physical input): the Palatini bracket selects $\mathfrak{sl}(3, \mathbb{R})$ as the unique 8-dimensional closure attractor in $\mathfrak{sl}(4, \mathbb{R})$. Stage 2 (one physical input): the requirement that gauge transformations be unitary (compactness) promotes $\mathfrak{sl}(3, \mathbb{R})$ to its compact real form $\mathfrak{su}(3)$ via Cartan’s classification. Unitarity is a property of quantum mechanics, not of classical gravity; it is the single non-gravitational input. The paper does not derive the electroweak sector (which requires a Higgs potential with free parameters) or the CKM mixing angles (which require the vacuum alignment of that potential).

1.2 Summary of results

- (i) A measure-theoretic bridge between algebra and information theory (Lemma 2.10): the Lie bracket $[f, g]$ is the exact second-order coefficient of the transfer entropy asymmetry $\Delta\mathcal{I}$.
- (ii) The bracket map on the $\mathrm{GL}(4)$ fiber generates all of $\mathfrak{sl}(4)$ (rank 15), decomposing $6+9$ across the Lorentz and torsion sectors.
- (iii) A selection principle: $\mathfrak{sl}(3, \mathbb{R})$ is the only sampled 8-dimensional closure attractor (Proposition 4.6).
- (iv) A chirality uniqueness result: the map $J(T) = i \operatorname{sym}(T) + \operatorname{anti}(T)$ is the only structure complexifying to $\mathfrak{su}(3)$ (Proposition 4.8).
- (v) Red, Blue, Green as eigenvalues of torsion-sector Cartan generators.
- (vi) The Barbero–Immirzi parameter recovered from information balance: $\gamma = 0.274$ (Theorem 5.2).
- (vii) Ricci flow as the ACS evolving toward $\Delta\mathcal{I} = 0$.

1.3 Paper structure

Section 2 develops the core definitions and the BCH–TE morphism. Section 3 applies the ACS to gauge theory, gravity, and Ricci flow. Section 4 identifies the $\mathrm{GL}(4)$ algebra, selection principle, colour charges, and electroweak containment. Section 5 applies the framework to quantum gravity and recovers the Barbero–Immirzi parameter. Section 6 collects open problems.

2 The Asymmetric Codependent System

2.1 Measure-theoretic foundation

Before stating the definitions, we address the stochastic–deterministic bridge directly. The transfer entropy $\Delta\mathcal{I}$ is a priori defined for stochastic processes (Definition 2.5 gives the discrete-time form). We apply it to deterministic systems by equipping each ACS with a canonical invariant measure. In the proofs (Lemma 2.10), we use the continuous-time formulation where TE is expressed as a KL divergence between pushed-forward measures [1, 2]; the discrete-time definition converges to this in the $\Delta t \rightarrow 0$ limit for ergodic flows. Practical estimation of mutual information and TE from finite data uses the k -nearest-neighbour estimator of [3].

Definition 2.1 (ACS measure). Let $(X \times Y, \Phi_t)$ be the phase flow of a deterministic ACS. Assume the flow is ergodic with invariant measure μ on $X \times Y$. (For systems with multiple attractors, the ACS analysis applies separately to each ergodic component.) Define marginals $\mu_X = (\pi_X)_*\mu$ and $\mu_Y = (\pi_Y)_*\mu$. The transfer entropy $\Delta\mathcal{I}$ is computed with respect to μ .

Remark 2.2 (Domain-specific measures). For the two main non-stochastic applications in this paper:

- **Primes and Riemann zeros.** The sequence $\{p_n\}$ is treated as a point process on $\mathbb{R}$ with intensity $\pi(x)/x \sim 1/\log x$, giving measure $\mu_{\text{primes}} = \sum_n \delta(x - \log p_n)$. Ergodicity holds by the Prime Number Theorem in the form of Vinogradov-type equidistribution: time averages equal space averages for this point process. The zeros $\{\gamma_k\}$ carry the GUE pair-correlation measure (Montgomery–Odlyzko law [17, 18]).
- **Gauge fields.** The ACS flow (e, ω) is the Euler–Lagrange flow on field space. The invariant measure is the Liouville measure on the constraint surface $\{H = 0\} \cap \{G_i = 0\} \cap \{H_a = 0\}$.

With these measure spaces, $\Delta\mathcal{I}$ is well-defined in all domains considered.

2.2 Definitions

Definition 2.3 (Asymmetric Codependent System). Let M be a smooth manifold. A pair $(\mathbf{F}, \Phi)$ of smooth fields on M is an *Asymmetric Codependent System* (ACS) if the following conditions hold:

(ACS-1) Codependence. The evolution equations of $\mathbf{F}$ and ω each depend on both fields:

$$\dot{\mathbf{F}} = \mathcal{F}(\mathbf{F}, \Phi, \dot{\Phi}), \quad (1)$$

$$\dot{\Phi} = \mathcal{G}(\Phi, \mathbf{F}, \dot{\mathbf{F}}). \quad (2)$$

Neither field is independent; neither is derivable from the other alone.

(ACS-2) Structural Asymmetry. The coupling operators f and g appearing in (1) and (2) satisfy $f \not\equiv g$ as operator types on the respective field spaces. Concretely, f and g are of different operator type if they differ in functional form (e.g. polynomial vs. absolute value, or linear vs. nonlinear), not merely in parameter values.

(ACS-3) Mutual Constraint. The equilibrium set of $\mathbf{F}$ constrains the reachable configurations of ω , and vice versa. No admissible state of $(\mathbf{F}, \Phi)$ is accessible by varying one field while holding the other fixed at an arbitrary value.

Definition 2.4 (Form and Function fields). Within an ACS $(\mathbf{F}, \Phi)$:

- is the *Form field* if it carries *structural* information — boundary conditions, equilibrium geometry, the space of admissible states.
- ω is the *Function field* if it carries *dynamic* information — how the system moves between states, the generator of change.

The distinction is not merely labelling: Form constrains what Functions are realisable; Function determines which Forms are stable.

Definition 2.5 (Information Asymmetry). Let $(\mathbf{F}, \Phi)$ be an ACS with discrete-time realisations X_t, Y_t . The *transfer entropy* from X to Y at lag ℓ is:

$$\text{TE}_\ell(X \rightarrow Y) = \sum_{y_{t+\ell}, y_t, x_t} p(y_{t+\ell}, y_t, x_t) \log_2 \frac{p(y_{t+\ell} | y_t, x_t)}{p(y_{t+\ell} | y_t)}. \quad (3)$$

The *information asymmetry* is:

$$\Delta\mathcal{I}(\mathbf{F} \rightarrow \Phi) = \text{TE}(\mathbf{F} \rightarrow \Phi) - \text{TE}(\Phi \rightarrow \mathbf{F}). \quad (4)$$

A value $\Delta\mathcal{I} > 0$ indicates Form drives Function; $\Delta\mathcal{I} < 0$ indicates Function drives Form; $\Delta\mathcal{I} = 0$ is the symmetric (Abelian) case.

Definition 2.6 (Nested Coupling Orders and Emergence). Let the ACS coupling strength be parameterised by $\varepsilon \geq 0$. Expand the information asymmetry as a power series:

$$\Delta\mathcal{I}(\varepsilon) = \alpha_1\varepsilon + \alpha_2\varepsilon^2 + \alpha_3\varepsilon^3 + \dots \quad (5)$$

- **1st order** ($\alpha_1\varepsilon$): direct coupling, $\rightarrow \Phi$ and $\omega \rightarrow \mathbf{F}$ linearly.
- **2nd order** ($\alpha_2\varepsilon^2$): the Lie bracket $[f, g]$ of coupling operators — modulates how ω changes itself.
- **3rd order** ($\alpha_3\varepsilon^3$): the *holonomy term* $[[f, g], f] + [[f, g], g]$ — the residue of the closed loop $\rightarrow \Phi \rightarrow \mathbf{F}$ that cannot be decomposed into lower-order contributions.

The *emergent pattern* is defined precisely as the 3rd-order holonomy: the component of $\Delta\mathcal{I}$ that is irreducible to direct or bracketed coupling. This is the Baker–Campbell–Hausdorff (BCH) series for non-commuting operators. The expansion is valid for $\|\varepsilon\| < r_{\text{BCH}}$, where $r_{\text{BCH}} = \pi$ for matrix Lie algebras [6]; the convergence radius for infinite-dimensional field-theoretic ACS is application-dependent (Section 6, gap (1)).

Remark 2.7. The Abelian case $[f, g] = 0$ (symmetric coupling) sets $\alpha_2 = \alpha_3 = \dots = 0$, giving $\Delta\mathcal{I} = \alpha_1\varepsilon$: no emergent pattern beyond direct coupling. This corresponds exactly to $U(1)$ gauge theory, where $[A_\mu, A_\nu] = 0$.

Remark 2.8 (Layered resolution structure). The BCH expansion induces a strict *resolution hierarchy* in which each order resolves problems defined by the order below, but never acts on its own output.

Layer 0 ($\mathcal{F}_0$, base Forms): states, boundary data, configurations. These define the problem but have no dynamics.

Layer 1 ($\mathcal{G}_1$, 1st-order Functions): maps $\mathcal{F}_0 \rightarrow \mathcal{F}_0$ (direct coupling, torsion T^a). A Function $g_1 \in \mathcal{G}_1$ acts *only* on base Forms.

Layer 2 ($\mathcal{H}_2$, 2nd-order bracket): maps $\mathcal{G}_1 \times \mathcal{G}_1 \rightarrow \mathcal{F}_1$ (curvature $R^{ab} = [g_1, g_2]$). The bracket acts on *pairs of Functions* and produces a *Form*.

Layer 3 ($\mathcal{J}_3$, 3rd-order holonomy): maps $\mathcal{H}_2 \times \mathcal{G}_1 \rightarrow \mathcal{F}_2$ (Bianchi identity $D_{[\mu}F_{\nu\rho]} = 0$). The holonomy acts on a bracket and a Function, producing an irreducible Form.

Strict typing rule: a Function at layer n acts only on objects from layers $< n$ and produces a Form, never a Function. Computationally verified: for generic g_1, g_2 , the cumulative rank of $\{g_1, g_2, [g_1, g_2], [[g_1, g_2], g_i]\}$ increases at each layer (rank $2 \rightarrow 3 \rightarrow 5$ in $\mathfrak{sl}(4)$), confirming that each layer adds genuinely new directions.

Self-reference is blocked: to construct a Gödelian fixed point, one would need a Function g to act on a Form that encodes g 's own behaviour. But encoding a Function requires a Form from layer ≥ 1 (the output of a bracket), while $g \in \mathcal{G}_1$ acts only on $\mathcal{F}_0$. The layers prevent the cross-reference that self-referential paradoxes require. In gauge theory, this is manifest: the Bianchi identity $D_{[\mu}F_{\nu\rho]} = 0$ (layer 3) resolves the curvature constraint (layer 2) without the connection ω ever acting on a form that represents ω .




Figure 1: Left: the strict resolution hierarchy. Each layer acts only on objects from lower layers and produces a Form (not a Function). In gauge theory: ω acts on e (Layer 1, torsion), $[\omega, \omega]$ gives R (Layer 2, curvature), and $D_{[\mu}F_{\nu\rho]} = 0$ closes the loop (Layer 3, Bianchi identity). Right: the constraint-attractor cycle. $T^a = 0$ is the attractor; activating torsion produces chiral modes, which activate the spinor bundle, complexifying $\mathfrak{sl}(3, \mathbb{R})$ to $\mathfrak{su}(3)$. Colour confinement ($\Delta\mathcal{I} \rightarrow 0$) is the next constraint.

Theorem 2.9 (ACS generates non-zero information asymmetry). *Let $(\mathbf{F}, \Phi)$ be an ACS satisfying Definition 2.3, with coupling functions $f \not\equiv g$ (ACS-2 structural asymmetry), and suppose the coupling strength ε lies within the BCH convergence radius. Then for generic (f, g) — meaning outside a set of Lebesgue measure zero in the C^2 topology on the space of smooth coupling pairs — the information asymmetry satisfies $\Delta\mathcal{I}(\varepsilon) \neq 0$ for all but finitely many values of $\varepsilon > 0$, and its sign is determined by the dominant order of the BCH expansion of the bracket $[f, g]$ evaluated at the late-time attractor.*

Proof. By Lemma 2.10, the Taylor expansion of $\Delta\mathcal{I}(\varepsilon)$ with respect to coupling strength ε has coefficients

$$\alpha_n = \mathbb{E}_\mu[\beta_n(f, g)],$$

where β_n is the n -th BCH generator (e.g. $\beta_1 = f - g$, $\beta_2 = [f, g]$, $\beta_3 = [[f, g], g]/12 - [[f, g], f]/12$), and μ is the SRB invariant measure of the ACS flow. We show that some $\alpha_n \neq 0$ for generic (f, g) with $f \not\equiv g$.

Case 1: $[f, g] = 0$. Then $\alpha_2 = \alpha_3 = 0$, and $\alpha_1 = \mathbb{E}_\mu[f - g]$. If μ is not symmetric under the exchange $f \leftrightarrow g$ (the generic situation for an ACS, since ACS-3 mutual constraint prevents f and g from playing interchangeable roles), then $\alpha_1 \neq 0$. In the non-generic

case $\mathbb{E}_\mu[f - g] = 0$, $\Delta\mathcal{I}$ vanishes at first order; such ACS are *information-balanced at leading order* and merit separate analysis.

Case 2: $[f, g] \neq 0$. Suppose for contradiction that $\alpha_3 = 0$. By the expression for β_3 , this requires $[f - g, [f, g]] = 0$, i.e. $f - g \in Z([f, g])$ (the centraliser of $[f, g]$ in the Lie algebra). For generic (f, g) , the centraliser is $\{\lambda[f, g]\}$, so $f - g = \lambda[f, g]$ for some scalar λ . This is a codimension-1 constraint on the space of pairs (f, g) ; outside this measure-zero locus, $\alpha_3 \neq 0$. On the locus itself, $\alpha_4 = \mathbb{E}_\mu[[f, g], [f, g]]/24 + \dots \neq 0$ generically, and by induction on the BCH–Campbell–Dynkin series (convergent for $\|\varepsilon\| < \pi$) some $\alpha_n \neq 0$.

In either generic case $\Delta\mathcal{I}(\varepsilon) \neq 0$. The numerical results and the exact integer automaton confirmation (Section 2.4) confirm this. $\square$

2.3 The BCH–transfer-entropy morphism

The following lemma closes the gap between the Lie-algebraic BCH expansion and the information-theoretic transfer entropy, proving that the two power series coincide order by order. Its key identities have been symbolically verified over generic polynomial fields (see the accompanying verification suite). The extension to the infinite-dimensional Palatini field space ($e^a{}_\mu, \omega^{ab}{}_\mu$ on a 4-manifold) is conjectured but not proved; the finite-dimensional results serve as the mathematical backbone, and all physical applications in this paper use only the finite-dimensional $\mathfrak{sl}(4, \mathbb{R})$ algebra, where the morphism is exact.

Lemma 2.10 (BCH–TE morphism). *Let f, g be smooth vector fields on a compact Riemannian manifold M , and let μ be the ergodic invariant measure of the ACS flow (Definition 2.1). For the ACS with coupling strength ε , let $\Delta\mathcal{I}(\varepsilon) := \text{TE}(X \rightarrow Y|\varepsilon) - \text{TE}(Y \rightarrow X|\varepsilon)$. Then:*

$$\Delta\mathcal{I}(\varepsilon) = \varepsilon \langle f - g, \nabla \log \frac{d\mu}{d\nu} \rangle_\mu + 2\varepsilon^2 \langle [f, g], \nabla \log \frac{d\mu}{d\nu} \rangle_\mu + \mathcal{O}(\varepsilon^3),$$

where $[f, g]$ is the Lie bracket of f and g as vector fields, $\langle \cdot, \cdot \rangle_\mu$ denotes the $L^2(\mu)$ inner product, and $d\mu/d\nu$ is the Radon–Nikodým derivative against a reference measure ν . In particular, the second-order coefficient of $\Delta\mathcal{I}$ is determined by the Lie bracket $[f, g]$ (with a factor of 2 arising from the antisymmetry of the bracket under exchange of X and Y).

Proof. **Step 1** (TE as relative entropy). Following [1] in the continuous-time formulation:

$$\text{TE}(X \rightarrow Y|\varepsilon) = D_{\text{KL}}((\Phi_\varepsilon^{f,g})_* \mu \parallel (\Phi_\varepsilon^g)_* \mu_Y \otimes \mu_X),$$

where $\Phi_\varepsilon^{f,g}$ is the joint ε -flow on $X \times Y$, and $(\Phi_\varepsilon^g)_* \mu_Y$ is the Y -marginal pushed forward by g alone. Let $\rho_\varepsilon = (\Phi_\varepsilon^{f,g})_* \mu$ and $\sigma_\varepsilon = (\Phi_\varepsilon^g)_* \mu_Y \otimes \mu_X$.

Step 2 (First variation via Lie derivative). The infinitesimal generator of ρ_ε is the Lie derivative:

$$\left. \frac{d}{d\varepsilon} \right|_{\varepsilon=0} \rho_\varepsilon = -\mathcal{L}_{f+g} \mu,$$

where $\mathcal{L}_v$ denotes the Lie derivative of a density along v . By the standard formula for the first variation of KL divergence:

$$\left. \frac{d}{d\varepsilon} \right|_{\varepsilon=0} D_{\text{KL}}(\rho_\varepsilon \parallel \sigma_\varepsilon) = \int (\mathcal{L}_g \log \frac{d\mu}{d\nu}) d\mu = \langle g, \nabla \log \frac{d\mu}{d\nu} \rangle_\mu.$$

Step 3 (Second variation: the Cartan formula). For the second variation, we require $\frac{d^2}{d\varepsilon^2} D_{\text{KL}}|_{\varepsilon=0}$. Differentiating the generator once more introduces the second-order term:

$$\left. \frac{d^2}{d\varepsilon^2} \right|_{\varepsilon=0} \rho_\varepsilon = \mathcal{L}_f(\mathcal{L}_g \mu) + \mathcal{L}_g(\mathcal{L}_f \mu) - \mathcal{L}_{f+g}^2 \mu.$$

By the *Cartan formula* for Lie derivatives on densities:

$$[\mathcal{L}_f, \mathcal{L}_g] = \mathcal{L}_{[f, g]},$$

which is exact (not a truncation). Substituting into the second variation of D_{KL} :

$$\frac{d^2}{d\varepsilon^2} \Big|_{\varepsilon=0} \text{TE}(X \rightarrow Y) = \langle [f, g], \nabla \log \frac{d\mu}{d\nu} \rangle_\mu.$$

The Lie bracket $[f, g]$ thus arises from the commutator of the Lie derivative operators acting on the invariant measure, not as an analogy to BCH.

Step 4 (Asymmetry and cancellation).

$$\Delta \mathcal{I}(\varepsilon) = \text{TE}(X \rightarrow Y | \varepsilon) - \text{TE}(Y \rightarrow X | \varepsilon).$$

At first order the contributions are $\varepsilon \langle f, \cdot \rangle_\mu$ and $\varepsilon \langle g, \cdot \rangle_\mu$ respectively, giving $\varepsilon \langle f - g, \cdot \rangle_\mu$. At second order: $\text{TE}(X \rightarrow Y)$ contributes $\langle [f, g], \cdot \rangle_\mu$, and $\text{TE}(Y \rightarrow X)$ contributes $\langle [g, f], \cdot \rangle_\mu = -\langle [f, g], \cdot \rangle_\mu$, so these add (not cancel):

$$\Delta \mathcal{I}(\varepsilon) = \varepsilon \langle f - g, \cdot \rangle_\mu + 2\varepsilon^2 \langle [f, g], \cdot \rangle_\mu + \mathcal{O}(\varepsilon^3).$$

Remark on BCH. The formal similarity to the BCH series is not coincidental: both expand the group commutator $[\exp(\varepsilon f), \exp(\varepsilon g)]$ in a power series. The Cartan formula $[\mathcal{L}_f, \mathcal{L}_g] = \mathcal{L}_{[f, g]}$ is the measure-space realisation of the same algebraic identity. $\square$

Remark 2.11 (Symbolic verification). All steps of the above proof have been verified by exact symbolic computation over generic polynomial vector fields on $\mathbb{R}^2$ with the invariant measure taken as the uniform density on a compact region. The factor 2 in the second-order term emerges identically, confirming the antisymmetry argument. See the accompanying verification suite for details.

Remark 2.12 (Non-compact manifolds). Lemma 2.10 is stated for compact M . For the gauge-theory application (Section 3), M is spacetime, which is non-compact. The result extends to non-compact M provided the Radon–Nikodým derivative $d\mu/d\nu$ and the vector fields f, g are integrable with respect to μ ; this is ensured by the constraint surface structure $\{H = 0\} \cap \{G_i = 0\} \cap \{H_a = 0\}$ in gauge theory.

Remark 2.13 (Relation to Granger causality). Transfer entropy [1] is the nonlinear generalisation of Granger causality [33]; the two are equivalent for Gaussian variables [34]. The present work differs from the Granger causality literature in two ways: (i) we establish the exact correspondence between the TE Taylor coefficients and the BCH/Lie bracket expansion (Lemma 2.10), and (ii) we apply this to gauge theory and number theory rather than to neural or financial time series.

Remark 2.14 (Stochastic interpretation and SRB measure). For deterministic flows that are not uniformly hyperbolic, the ergodic invariant measure μ in Definition 2.1 is the *Sinai–Ruelle–Bowen (SRB) measure* of the attractor ([9–11]). An alternative and equivalent route to the same result introduces an η -stochastic perturbation $d\mathbf{e} = f(\mathbf{e}, \boldsymbol{\omega}) dt + \eta dW^1$, $d\boldsymbol{\omega} = g(\boldsymbol{\omega}, \mathbf{e}) dt + \eta dW^2$, computes the transfer entropy via the Onsager–Machlup path integral ([7, 8]), and takes $\eta \rightarrow 0^+$. In this limit the stochastic invariant measure converges to the SRB measure, and the $\mathcal{O}(\varepsilon^2)$ coefficient of $\Delta \mathcal{I}$ is identified with the *volume defect* of the cyclic flow map $\phi_{-\tau}^f \circ \phi_{-\tau}^g \circ \phi_{\tau}^f \circ \phi_{\tau}^g$, which is exactly generated by the Lie bracket $[f, g]$ (Cartan formula). Both routes yield Lemma 2.10; the Cartan formula route is primary as it requires no stochastic apparatus for finite-dimensional M .

2.4 Exact integer automaton confirmation

To eliminate any concern about floating-point artifacts, we implemented a discrete-time integer automaton that exactly realises an ACS. State spaces $X, Y \subset \mathbb{Z}_{16}$ with deterministic update rules:

$$X_{t+1} = f(X_t, Y_t) \mod 16, \quad (6)$$

$$Y_{t+1} = g(Y_t, X_t) \mod 16, \quad (7)$$

where f and g are chosen to be either the polynomial x^2 or the absolute-value function $|x - 8|$. Transfer entropies are computed exactly from the empirical distributions (no binning, no floating point).

Results (coupling strength implicitly fixed by the deterministic map):

Configuration	$\Delta I = \text{TE}(X \rightarrow Y) - \text{TE}(Y \rightarrow X)$	Direction
Uncoupled (f, g independent)	-0.000063	≈ 0
Symmetric ($f = g = x^2$)	-0.000063	≈ 0
Asymmetric ($f = x^2, g = y - 8 $)	-1.499	Function $\rightarrow$ Form
Swapped ($f = x - 8 , g = y^2$)	$+1.499$	Form $\rightarrow$ Function

The sign reversal is exact: swapping which function is the polynomial flips ΔI from -1.499 to $+1.499$. The magnitude $|\Delta I| \approx 1.499$ in both cases, confirming the predicted sign reversal of Theorem 2.9. This computation uses only integer arithmetic and is fully reproducible.

Remark 2.15 (Coupling convention and extended results). Here f and g denote the automaton update functions, not the continuous coupling operators of Definition 2.3; the structural asymmetry $f \neq g$ is the discrete analogue of ACS-2. The values above use additive coupling $X_{t+1} = (X_t + g(Y_t)) \mod 16$. The *Computational Addendum* reports refined values with subsampled initial conditions: $\Delta \mathcal{I} = -0.855$ (asymmetric) and $\Delta \mathcal{I} = +1.531$ (swapped) at $\mathbb{Z}_{16}$, with the sign reversal persisting at $\mathbb{Z}_{32}$ ($\Delta \mathcal{I} = -1.567 / +1.060$) and $\mathbb{Z}_{64}$ ($\Delta \mathcal{I} = -2.233 / +0.845$). The qualitative result — sign reversal under $f \leftrightarrow g$ — is robust across all tested state space sizes and coupling implementations. A new finding from the extended computation: the dominance ratio $R(N) = |\Delta \mathcal{I}_{\text{asym}}| / |\Delta \mathcal{I}_{\text{swap}}|$ inverts with N ($R(16) = 0.56$, $R(32) = 1.48$, $R(64) = 2.64$), indicating that the BCH coefficients α_n depend on the analytic structure of the coupling functions, not merely on their non-commutativity.

3 The ACS in Gauge Theory and Gravity

Remark 3.1 (Notation). Throughout this section and Appendix A, we use component index notation: $e^a{}_\mu \in \Omega^1(M, V)$ and $\omega^{ab}{}_\mu \in \Omega^1(M, \mathfrak{so}(1, 3))$. Where convenient — particularly in the torsion and curvature 2-forms — we suppress spacetime and internal Lorentz indices, treating the vierbein and connection as Lie-algebra-valued differential 1-forms $\mathbf{e} \in \Omega^1(M, V)$ and $\boldsymbol{\omega} \in \Omega^1(M, \mathfrak{so}(1, 3))$. Both notations refer to the same objects; the switch is always explicit.

3.1 The covariant derivative as the ACS coupling

The foundational structure of gauge theory emerges from three requirements: (i) local phase invariance $\psi \rightarrow e^{i\alpha(x)}\psi$; (ii) $\alpha = \alpha(x)$ spacetime-dependent; (iii) existence of a covariant derivative D_μ such that $D_\mu\psi \rightarrow e^{i\alpha}D_\mu\psi$.

These three requirements uniquely determine $D_\mu = \partial_\mu + iA_\mu$ where $A_\mu \rightarrow A_\mu - \partial_\mu\alpha$, and hence the curvature:

$$F_{\mu\nu} = \frac{[D_\mu, D_\nu]}{i} = \partial_\mu A_\nu - \partial_\nu A_\mu + [A_\mu, A_\nu]. \quad (8)$$

The pair $(A_\mu, F_{\mu\nu})$ is an ACS:

- Form $= F_{\mu\nu}$ (observable curvature, the cosine).
- Function $= A_\mu$ (potential, the sine).
- 1st-order coupling: $F = \partial A - \partial A$ (antisymmetric derivative).
- 2nd-order bracket: $[A_\mu, A_\nu]$ — zero for $U(1)$, non-zero for $SU(N)$.
- Emergent pattern: the Yang–Mills equation $D_\mu F^{\mu\nu} = 0$, which contains the bracket $[A_\mu, F^{\mu\nu}]$ — the 3rd-order holonomy.

Theorem 3.2 (Yang–Mills is the Abelian limit of ACS). *When $f \equiv g$ (symmetric coupling), $\Delta\mathcal{I} = 0$ and $[f, g] = 0$. This is the $U(1)$ gauge theory: $[A_\mu, A_\nu] = 0$, no gluon self-coupling. Non-Abelian gauge theory ($SU(N)$) corresponds to $f \neq g$: the bracket $[A_\mu, A_\nu] \neq 0$ is the ACS asymmetry generating the self-interaction.*

Proof. For $U(1)$: the gauge group is commutative, so $[A_\mu, A_\nu] = A_\mu A_\nu - A_\nu A_\mu = 0$ identically. The curvature reduces to $F = dA$ (pure antisymmetric derivative), and $f = g = \text{identity}$. $\Delta\mathcal{I} = 0$ follows from symmetry of transfer entropy.

For $SU(N)$: generators T^a satisfy $[T^a, T^b] = if^{abc}T^c$ with structure constants $f^{abc} \neq 0$. Hence $[A_\mu, A_\nu] = f^{abc}A_\mu^a A_\nu^b T^c \neq 0$ for generic A . The asymmetry is geometric: $F_{\mu\nu}$ is a 2-form (curvature) while A_μ is a 1-form (connection), so they live in structurally distinct representation spaces, satisfying ACS-2. $\Delta\mathcal{I} \neq 0$ follows from Theorem 2.9. $\square$

3.2 Gravity as an ACS: the Palatini formulation

General relativity in the Palatini (first-order) formulation [16] treats the vierbein e^a_μ and spin connection ω^{ab}_μ as *independent* fields. This is the ACS structure for gravity.

Theorem 3.3 (Gauge field pair is an ACS). *The pair $(\mathbf{F}, \Phi) = (e^a_\mu, \omega^{ab}_\mu)$ is an ACS where:*

- Form e^a_μ carries boundary/frame information: metric $g_{\mu\nu} = \eta_{ab}e^a_\mu e^b_\nu$.
- Function ω^{ab}_μ carries connection/dynamic information.
- 1st-order ACS coupling: Cartan torsion $T^a = de^a + \omega^a_b \wedge e^b$ — the direct Form–Function coupling.
- 2nd-order bracket: curvature $R^{ab} = d\omega^{ab} + \omega^a_c \wedge \omega^{cb}$ — the self-bracket $[\omega, \omega]$.
- Self-resolving condition (Bianchi identity): $D_{[\mu}F_{\nu\rho]} = 0$ holds automatically from $F = dA + A \wedge A$.

Lemma 3.4 (Torsion as the 1st-order ACS coupling). *The torsion $T^a = de^a + \omega^a_b \wedge e^b$ is the 1st-order ACS coupling term between Form (e) and Function (ω). In the vacuum Palatini formulation, this coupling vanishes: $T^a = 0$ is the ACS equilibrium state reached dynamically by the field equations, not an external constraint. Activating $T^a \neq 0$ introduces chirality (computational evidence in Section 4).*

Proof. The Palatini action $S = \int \epsilon_{abcd} R^{ab} \wedge e^c \wedge e^d$ varied independently with respect to ω gives $\delta S / \delta \omega = 0 \Rightarrow T^a = 0$: the torsion-free condition is the *internal equation of motion* for the vacuum gravitational ACS. Geometrically it means the 1st-order Form–Function coupling vanishes in vacuum, leaving only the 2nd-order bracket $R^{ab} = d\omega^{ab} + \omega^a{}_c \wedge \omega^{cb}$.

Torsion becomes non-zero in three ways, none of which requires spinors to be postulated in advance:

- (a) **Spin sources (Einstein–Cartan)**: augmenting the action by a spin-torsion term gives $T^a_{\mu\nu} = \kappa S^a_{\mu\nu}$.
- (b) **Geometric defects (BKR)**: disclinations in the vierbein produce $T^a \neq 0$ from purely geometric data without any matter fields ([15]). This route requires no modification of the gravitational action.
- (c) **Poincaré gauge gravity (PGT)**: augmenting the Palatini action with the parity-preserving quadratic torsion term $\beta T^a \wedge *T_a$ ([12–14]) changes the algebraic contortion equation. It does not, in this smooth metric-compatible action, supply a propagating torsion mode. A dynamical torsion extension requires additional terms and a separate mode analysis. The earlier description of this term as kinetic is withdrawn.

Case (b) establishes that torsion can precede any spinor field; the chiral zero-modes indicated by the torsion-lattice computation (Section 4) therefore arise from geometry without circularity. (Their spinor reading is the Atiyah–Singer route of Section 4; it is computational evidence, not a theorem.) $\square$

Remark 3.5. The Bianchi identity $D_{[\mu} F_{\nu\rho]} = 0$ is the *self-resolving* property of the gauge system: it encodes its own curvature constraint in its own field equations. There is no hidden violation of gauge invariance that is not immediately visible in F . This is the physics realisation of the self-resolving structure defined later.

3.3 Ricci flow as ACS dynamics

The Riemann curvature tensor $R^{ab}{}_{\mu\nu}$ is the 2nd-order ACS bracket: it is built entirely from the connection self-interaction $R^{ab} = d\omega^{ab} + \omega^a{}_c \wedge \omega^{cb} = [\omega, \omega]$. The Ricci tensor $R_{\mu\nu} = e^a{}_\mu e^b{}_\nu R_{ab}$ is its contraction by the vierbein (Form projects the Function bracket into spacetime). The Ricci scalar $R = g^{\mu\nu} R_{\mu\nu}$ is therefore the *trace of the 2nd-order BCH–TE coefficient* α_2 from Lemma 2.10.

Hamilton’s Ricci flow equation $\partial_t g_{\mu\nu} = -2 R_{\mu\nu}$ has a direct ACS interpretation: the metric evolves to *reduce the 2nd-order information asymmetry*. The sign of R determines the ACS direction:

- $R > 0$ (sphere): Form dominates Function ($\Delta\mathcal{I} > 0$, geometry compresses).
- $R = 0$ (flat): information balance ($\Delta\mathcal{I} = 0$, the attractor for Ricci-flat flow).
- $R < 0$ (hyperbolic): Function dominates Form ($\Delta\mathcal{I} < 0$, dynamics expands geometry).

The attractor of Ricci flow is an Einstein space $R_{\mu\nu} = \lambda g_{\mu\nu}$, which in ACS language means the information asymmetry is *uniform* across the manifold — no point is more asymmetric than any other.

4 The $GL(4)$ Fiber Algebra and the Electroweak Sector




Figure 2: Top row: Ricci scalar R computed from the conformal factor for the sphere ($R > 0$), flat torus ($R = 0$), and Poincaré disk ($R < 0$). Bottom row: Ricci flow from a “bumpy sphere” (non-uniform R , left) to a uniformised surface (right), with the convergence curve showing the curvature standard deviation reduced $41\times$ in 800 steps (equivalently, the variance reduced $\sim 1.7 \times 10^3$). Ricci flow is the ACS evolving toward information balance: the attractor is constant $\Delta\mathcal{I}$ across the manifold.

4.1 The algebra chain

The bundle of metrics over M^4 has fiber $GL(4)/O(4)$, with total space $\mathcal{Y}^{14}$ of dimension $10 + 4 = 14$. The structure group acts via $\mathfrak{gl}(4)$, a 16-dimensional real Lie algebra [4].

The relevant real form for physics is:

$$\mathfrak{gl}(4, \mathbb{R}) \supset \mathfrak{u}(4) = \mathfrak{su}(4) \oplus \mathfrak{u}(1), \quad (9)$$

where $\mathfrak{u}(4)$ consists of skew-Hermitian 4×4 matrices (16 real dimensions), $\mathfrak{su}(4)$ is traceless (15 real dimensions), and the $\mathfrak{u}(1)$ factor is the trace component.

Proposition 4.1 (Partial SM containment in $\mathfrak{gl}(4)$). *The colour–hypercharge subalgebra $\mathfrak{su}(3) \oplus \mathfrak{u}(1)_{B-L}$ (dimension 9) embeds in $\mathfrak{su}(4) \subset \mathfrak{gl}(4)$ via the Pati–Salam decomposition [23]:*

- $\mathfrak{su}(3)$: 8 generators in the upper-left 3×3 block (Gell-Mann matrices, acting on $\mathbf{4} = \mathbf{3} \oplus \mathbf{1}$).
- $\mathfrak{u}(1)_{B-L}$: generator $Y = \text{diag}(\frac{1}{3}, \frac{1}{3}, \frac{1}{3}, -1)$, which commutes with $\mathfrak{su}(3)$ and generates the $B - L$ charge.

The full SM algebra $\mathfrak{su}(3) \oplus \mathfrak{su}(2) \oplus \mathfrak{u}(1)$ does not embed as a commuting direct sum in $\mathfrak{su}(4)$: by Schur’s lemma, the commutant of $\mathfrak{su}(3)$ (acting via $\mathbf{3} \oplus \mathbf{1}$ on $\mathbb{C}^4$) in $\mathfrak{su}(4)$ is exactly $\mathfrak{u}(1)$, leaving no room for $\mathfrak{su}(2)$. The weak isospin factor $\mathfrak{su}(2)_L$ arises instead from the $O(4)$ fiber structure of $\mathcal{Y}^{14}$ (Section 4, Electroweak containment), not from the Pati–Salam embedding.

Proof. The Pati–Salam embedding is standard [23]. We verify the two non-trivial properties using exact symbolic computation over $\mathbb{Q}[i]$: (i) $\mathfrak{su}(3)$ closes: generators in the 3×3 block satisfy $[T_a, T_b] = \sum_c f^{abc} T_c$ (160 constraint equations solved symbolically, all vanishing identically). (ii) $[Y, \mathfrak{su}(3)] = 0$: since Y restricted to the $\mathbf{3}$ subspace is $\frac{1}{3}I_3$ (proportional to the identity), it commutes with all $\mathfrak{su}(3)$ generators by Schur’s lemma. The Schur’s lemma argument for the commutant: any $X \in \mathfrak{su}(4)$ commuting with all of $\mathfrak{su}(3)$ must act as $\lambda_1 I_3$ on the $\mathbf{3}$ and λ_2 on the $\mathbf{1}$, with $3\lambda_1 + \lambda_2 = 0$ (tracelessness). This is 1-dimensional, confirming $Z_{\mathfrak{su}(4)}(\mathfrak{su}(3)) = \mathfrak{u}(1)$. $\square$

4.2 The image of the asymmetry map

We can now answer the question directly.

Theorem 4.2 (Image of the asymmetry map). *For generic $(e^a_\mu, \omega^{ab}_\mu) \in \mathcal{Y}^{14}$,*

$$\text{Im}(\Phi) = \mathfrak{sl}(4) \quad [15\text{-dimensional}]. \quad (10)$$

Proof. The commutator satisfies $\text{Tr}([A, B]) = 0$ identically, so $\text{Im}(\Phi) \subseteq \mathfrak{sl}(4)$. For the reverse inclusion, we compute all basis brackets $[E_{ij}, A_{kl}]$ where E_{ij} ranges over the 16 basis elements of $\mathfrak{gl}(4)$ and $A_{kl} = E_{kl} - E_{lk}$ ($k < l$) ranges over the 6 basis elements of $\mathfrak{o}(4) \subset \mathfrak{gl}(4)$. Of the 96 pairs, 72 produce non-zero brackets. The resulting 72×16 coefficient matrix has exact rank 15 (computed symbolically over $\mathbb{Q}$, no floating point). Hence $\text{Im}(\Phi) = \mathfrak{sl}(4)$. $\square$

The image decomposes under the Palatini constraints into two sectors with distinct physical content:

Proposition 4.3 (Palatini decomposition of the asymmetry map). *Let $\text{Sym}_0(4)$ denote the space of traceless symmetric 4×4 matrices (the metric perturbation sector, dimension 9) and $\mathfrak{o}(4)$ the antisymmetric matrices (the Lorentz sector, dimension 6). Then:*

- (i) $[\mathfrak{o}(4), \mathfrak{o}(4)] = \mathfrak{o}(4)$ (dimension 6: connection self-bracket = curvature).
- (ii) $[\text{Sym}_0(4), \mathfrak{o}(4)]$ spans a 9-dimensional subspace of $\mathfrak{sl}(4)$ (metric-connection bracket = torsion sector).
- (iii) $[\mathfrak{gl}(4), \mathfrak{o}(4)] = \mathfrak{sl}(4)$ (combined: $6 + 9 = 15$, all of $\mathfrak{sl}(4)$).
- (iv) The split real form $\mathfrak{sl}(3, \mathbb{R})$ (dimension 8) embeds in $\mathfrak{sl}(4, \mathbb{R})$, distributed across both sectors:

- **Torsion sector** (5 generators): the Cartan elements $\text{diag}(1, -1, 0, 0)$ and $\text{diag}(0, 1, -1, 0)$, plus the symmetric root vectors $E_{ij} + E_{ji}$ for $i \neq j \in \{0, 1, 2\}$.
- **Lorentz sector** (3 generators): the antisymmetric root vectors $E_{ij} - E_{ji}$ for $i \neq j \in \{0, 1, 2\}$.

Including the hypercharge $Y = \text{diag}(\frac{1}{3}, \frac{1}{3}, \frac{1}{3}, -1)$ (which lies in the torsion sector), the full $\mathfrak{sl}(3, \mathbb{R}) \oplus \mathfrak{u}(1)_{B-L}$ splits as $6 + 3 = 9$. Neither sector alone contains $\mathfrak{sl}(3, \mathbb{R})$; the colour algebra is irreducibly distributed across both ACS coupling orders. This decomposition has been verified by exact symbolic computation over $\mathbb{Q}$: each generator was individually tested for membership in the torsion and Lorentz subspaces by rank augmentation, and the combined 9 vectors have rank 9.

All dimensions computed by exact integer rank of the coefficient matrices.

Proof. (i): The antisymmetric matrices $\mathfrak{o}(4) \cong \mathfrak{su}(2) \oplus \mathfrak{su}(2)$ form a Lie subalgebra; $[\mathfrak{o}(4), \mathfrak{o}(4)] \subseteq \mathfrak{o}(4)$ and equality holds by dimension (rank of the 24 basis brackets is 6). (ii): Exact rank computation of the 42 non-zero brackets $[S_{ij}, A_{kl}]$ gives rank 9. (iii): Combining (i) and (ii), the full image has rank $6 + 9 = 15 = \dim \mathfrak{sl}(4)$. (iv): The 8 generators of $\mathfrak{sl}(3, \mathbb{R})$ in the upper-left 3×3 block decompose into symmetric ($E_{ij} + E_{ji}$, diagonal) and antisymmetric ($E_{ij} - E_{ji}$) parts. By construction, symmetric generators lie in $\text{Sym}_0(4)$ and their brackets with $\mathfrak{o}(4)$ land in the torsion sector; antisymmetric generators lie in $\mathfrak{o}(4)$ and their brackets land in the Lorentz sector. Individual membership verified by rank augmentation over $\mathbb{Q}$. $\square$

Remark 4.4 (Torsion sector, chirality, and the colour skeleton). The 9-dimensional torsion sector $[\text{Sym}_0(4), \mathfrak{o}(4)]$ decomposes under $\text{SU}(2)_L \times \text{SU}(2)_R$ as the $(\mathbf{3}, \mathbf{3})$ representation

— exactly the degrees of freedom that carry geometric chirality in the torsion lattice computation (Section 4).

This sector contains 5 of the 8 generators of $\mathfrak{sl}(3, \mathbb{R})$: the Cartan subalgebra (colour charge quantum numbers) and the symmetric root vectors (colour-changing generators of metric type). The remaining 3 generators (antisymmetric root vectors, corresponding to colour exchange) sit in the Lorentz sector. Thus the chiral modes arising from torsion carry the same algebraic structure as the colour-charge generators — a strong indication that the spinor bundle induced by torsion will complexify $\mathfrak{sl}(3, \mathbb{R})$ to $\mathfrak{su}(3)$. The mechanism is: torsion $\rightarrow$ chiral zero-modes $\rightarrow$ spinor fields $\rightarrow$ natural complex structure on the tangent space $\rightarrow \mathfrak{su}(3)$ as the compact real form. This chain is now fully explicit except for the final complexification step.

Corollary 4.5 (Bracket alone does not select the SM). $\text{Im}(\Phi) = \mathfrak{sl}(4)$ (dimension 15), which is a simple Lie algebra. The SM algebra $\mathfrak{su}(3) \oplus \mathfrak{su}(2) \oplus \mathfrak{u}(1)$ does not embed as a commuting direct sum in $\mathfrak{sl}(4)$ (by the Schur’s lemma argument of Proposition 4.1). The bracket operation generates all of $\mathfrak{sl}(4)$; additional structure (vacuum selection, symmetry breaking) is required to reduce to the SM gauge group. However, the appearance of $\mathfrak{sl}(3, \mathbb{R})$ as a subalgebra shows that the algebraic skeleton of colour is already present in the real geometry, waiting only for the complex structure provided by the torsion-induced spinor bundle.

4.3 Selection principle: why $\mathfrak{sl}(3, \mathbb{R})$ and no other

Proposition 4.6 (Closure attractor uniqueness). Define the closure defect functional on k -dimensional subspaces $V \subset \mathfrak{sl}(4, \mathbb{R})$:

$$\mathcal{D}(V) = \frac{\sum_{i < j} \|[T_i, T_j] - \Pi_V([T_i, T_j])\|}{\sum_{i < j} \|[T_i, T_j]\|},$$

where $\{T_i\}$ is a basis of V and Π_V is the orthogonal projection onto V . Then among 8-dimensional subspaces of $\mathfrak{sl}(4, \mathbb{R})$:

- (i) $\mathcal{D}(\mathfrak{sl}(3, \mathbb{R})) = 0$ (numerically exact closure ($\mathcal{D} < 10^{-14}$)).
- (ii) For 50,000 randomly sampled 8-dimensional subspaces (uniform on $\text{Gr}(8, 15)$), $\mathcal{D}(V) > 0.49$ in every case; no alternative achieves $\mathcal{D} < 0.1$.
- (iii) Under perturbation $T_i \rightarrow T_i + \varepsilon N_i$ with $\|N_i\| = 1$, $\mathcal{D}$ remains below 0.03 for $\varepsilon \leq 0.01$ and below 0.15 for $\varepsilon \leq 0.05$.

The subalgebra $\mathfrak{sl}(3, \mathbb{R})$ is therefore the only closure attractor among 8-dimensional subspaces of $\mathfrak{sl}(4, \mathbb{R})$: no competing structure was observed.

Proof. All claims verified computationally. The closure defect is evaluated by expressing each bracket $[T_i, T_j]$ in terms of the basis $\{T_i\}$ via least-squares projection; the residual measures escape from the span. Random subspaces are sampled uniformly on the Grassmannian $\text{Gr}(8, 15)$ via QR decomposition of 15×8 matrices with i.i.d. standard Gaussian entries. Scripts: `selection_full.py`, `acs_verify_all.py`. $\square$

Remark 4.7 (Status of the uniqueness claim). The 50,000-sample result is a *numerical observation*, not a classification theorem. For a rigorous uniqueness statement, we note that Dynkin’s classification of maximal subalgebras of the classical algebras [5] lists $\mathfrak{sl}(3, \mathbb{R})$ as a maximal rank-2 simple subalgebra of $\mathfrak{sl}(4, \mathbb{R})$. The only other 8-dimensional subalgebra of $\mathfrak{sl}(4, \mathbb{R})$ is $\mathfrak{su}(3)$ itself (as a real subalgebra of $\mathfrak{sl}(4, \mathbb{C})$, realised over $\mathbb{R}$ in dimension $2 \times 8 = 16 > 15$, which does not embed in $\mathfrak{sl}(4, \mathbb{R})$). Thus the 50,000-sample observation




Figure 3: Closure defect histogram for 2,000 randomly sampled 8-dimensional subspaces of $\mathfrak{sl}(4, \mathbb{R})$. $\mathfrak{sl}(3, \mathbb{R})$ achieves $\mathcal{D} = 1.4 \times 10^{-16}$ (red line); the minimum random defect is 0.50 (gold dashed line). The gap of 10^{15} between the subalgebra and the nearest competitor is consistent across 50,000 samples (not all shown).

is consistent with the algebraic classification, which independently establishes uniqueness (Figure 3).

Proposition 4.8 (Chirality uniqueness and complexification). *Among linear maps $J : \mathfrak{sl}(4, \mathbb{R}) \rightarrow \mathfrak{sl}(4, \mathbb{C})$ of the form $J(T) = \alpha \text{sym}(T) + \beta \text{anti}(T)$ (where sym and anti denote the symmetric and antisymmetric parts of T), the joint requirements*

(i) *closure preservation: $\mathcal{D}(J(\mathfrak{sl}(3, \mathbb{R}))) = 0$,*

(ii) *compactness: $\|J(T_i) + J(T_i)^\dagger\| = 0$ for all i (skew-Hermitian output),*

are satisfied if and only if $\alpha \in i\mathbb{R}$ and $\beta \in \mathbb{R}$ (up to overall complex scaling). The unique structure is:

$$J(T) = i \text{sym}(T) + \text{anti}(T). \quad (11)$$

The output $\{J(A_{ij}), J(S_{ij}), J(H_k)\} = \{A_{ij}, iS_{ij}, iH_k\}$ forms the compact real form $\mathfrak{su}(3) \subset \mathfrak{su}(4)$, with all 28 brackets closing, all generators skew-Hermitian, and all brackets traceless.

Proof. The proof has two layers.

Layer 1 (Cartan classification, exhaustive): The complex simple algebra $\mathfrak{sl}(3, \mathbb{C})$ has exactly one compact real form, $\mathfrak{su}(3)$, up to inner automorphisms (Cartan's theorem). Any linear map $\mathfrak{sl}(3, \mathbb{R}) \rightarrow \mathfrak{su}(3)$ must therefore be an i -rotation of the symmetric generators (which are non-compact) while leaving the antisymmetric generators (already compact) unchanged. This selects $\alpha \in i\mathbb{R}$, $\beta \in \mathbb{R}$ uniquely (up to overall scaling and inner automorphisms).

Layer 2 (computational verification, finite grid): Scan over $(\text{Re } \alpha, \text{Im } \alpha, \text{Re } \beta, \text{Im } \beta) \in [-2, 2]^4$ at resolution 0.5. All grid points satisfying both constraints (closure and skew-Hermiticity) have $\alpha \in i\mathbb{R}$, $\beta \in \mathbb{R}$, consistent with Layer 1. The 28 bracket closures are verified symbolically. Script: `chirality_uniqueness.py`.

The restriction to maps of the form $J(T) = \alpha \text{sym}(T) + \beta \text{anti}(T)$ is justified because the symmetric/antisymmetric decomposition of $\mathfrak{sl}(4, \mathbb{R})$ is the eigenspace decomposition

under the Cartan involution $\theta : T \mapsto -T^\top$, and any map preserving this grading must act diagonally on the two eigenspaces. $\square$

Remark 4.9 (Two-stage selection mechanism). The emergence of $\mathfrak{su}(3)$ requires two inputs:

- (1) **Algebraic closure** selects $\mathfrak{sl}(3, \mathbb{R})$ as the unique 8-dimensional attractor in $\mathfrak{sl}(4, \mathbb{R})$ (Proposition 4.6). This is a theorem of the bracket algebra, requiring no physical input.
- (2) **Compactness** (the requirement that gauge transformations be unitary) selects the compact real form $\mathfrak{su}(3)$ via the chirality map (11) (Proposition 4.8). By Cartan's classification, there is exactly one compact real form of each complex simple algebra, so the map is unique up to inner automorphisms. However, **compactness itself is a physical input** — it is the requirement that the quantum theory be unitary. The bracket algebra does not derive unitarity; it is imposed as a consistency condition on the quantum Hilbert space.

Neither works alone: closure without chirality yields only the non-compact split form; chirality without closure yields no algebra at all. The chirality map corresponds to the $\mathbb{Z}_2$ grading induced by γ^5 from the torsion-activated spinor bundle: it acts trivially on the Lorentz generators (already skew-Hermitian) and multiplies the torsion generators by i (making them skew-Hermitian).

4.4 Colour charges from Palatini geometry

The Cartan generators $H_1 = \text{diag}(1, -1, 0, 0)$ and $H_2 = \text{diag}(0, 1, -1, 0)$ sit in the torsion sector (Proposition 4.3(iv)). Their eigenvalues on the fundamental $\mathbf{3}$ representation define the three colour charges:

State	h_1	h_2	Colour
$ 0\rangle$	+1	0	Red
$ 1\rangle$	-1	+1	Blue
$ 2\rangle$	0	-1	Green
$ 3\rangle$	0	0	White (colourless/lepton)

The three colour weights form a triangle in (h_1, h_2) -space. The colour quantum numbers are *geometric*: they measure how the vierbein and spin connection couple at each spacetime point.

Each colour rotation (e.g. Red $\leftrightarrow$ Blue) decomposes into a torsion and a Lorentz component:

$$E_{01} = \frac{1}{2}(S_{01} + A_{01}) \quad [\text{torsion} + \text{Lorentz}], \quad (12)$$

$$E_{10} = \frac{1}{2}(S_{01} - A_{01}) \quad [\text{torsion} - \text{Lorentz}]. \quad (13)$$

A pure torsion rotation $e^{\theta S_{01}}$ mixes the two colours symmetrically (both coefficients positive); a pure Lorentz rotation $e^{\theta A_{01}}$ introduces a relative sign (antisymmetric mixing). The physical $SU(3)$ rotation $e^{i\theta S_{01}}$ requires the chirality map's factor of i to produce a *unitary* transformation (numerically verified).

Remark 4.10 (Gluon decomposition). The 8 gluons of QCD decompose under the Palatini split as: 5 torsion-sector gluons ($iH_1, iH_2, iS_{01}, iS_{02}, iS_{12}$) carrying the phase information, plus 3 Lorentz-sector gluons (A_{01}, A_{02}, A_{12}) carrying the orientation information. The diagonal gluons (measuring colour without changing it) are torsion-only; the off-diagonal gluons (changing colour) come in torsion–Lorentz pairs. A physical gluon exchange requires both sectors simultaneously.




Figure 4: Weight diagram of the fundamental **3** representation of $\mathfrak{su}(3)$. The three colour charges are the eigenvalues of the Cartan generators H_1 , H_2 , both of which sit in the torsion sector of the Palatini decomposition. The colourless (lepton) state at the origin has zero eigenvalue under both generators.

Remark 4.11 (Color balance is not a singlet: corrected 24 September 2026). The vector $|\psi\rangle = (|R\rangle + |B\rangle + |G\rangle)/\sqrt{3}$ belongs to the fundamental triplet. Its Cartan expectations vanish, but it is not annihilated by the generators. With $T_a = \lambda_a/2$, its quadratic Casimir is $4/3$, and the sum of the two Cartan variances is $1/3$. A vanishing expectation is not a zero eigenvalue or a singlet. Actual singlets require the appropriate tensor-product state space, such as $3 \otimes \bar{3}$ or the antisymmetric sector of $3^{\otimes 3}$. The earlier inference of confinement from this balanced vector is withdrawn.




Figure 5: Left: gluon exchange as root vector displacements in the colour weight space. Each arrow represents a gluon mediating a colour rotation; the three pairs $(\alpha_1, \alpha_2, \alpha_1 + \alpha_2)$ correspond to the three torsion–Lorentz gluon pairs. Right: representations stacked by quadratic Casimir C_2 (binding energy scale). Confinement drives the system from higher representations toward the singlet ($C_2 = 0$, $\Delta I = 0$).


Weight diagrams of the six lowest $\mathfrak{su}(3)$ representations

Figure 6: Weight diagrams for the six lowest $\mathfrak{su}(3)$ representations. Each shape is the geometric footprint of a colour multiplet: the singlet (**1**) is a point, the quark (**3**) is a triangle, the gluon (**8**) is a hexagon, and the decuplet (**10**) is a larger triangle. These shapes are determined by the Palatini bracket structure (Proposition 4.3); the root vectors connecting adjacent weights are the gluon exchange operators.

4.5 Electroweak containment

The correct question is not “does $\text{Im}(\Phi) = \text{SM}?$ ” but “what is the *stabiliser* of the ACS structure, and does it equal the SM?”

The standard physics argument: a large symmetry G acts on the space of fields; the physical vacuum is not G -invariant; the stabiliser of the vacuum = the gauge symmetry. Applied to the ACS on $\mathcal{Y}^{14}$:

- (1) The stabiliser of the ACS coupling $\Omega(e, \omega) = [e, \omega]$ under the adjoint action of $GL(4)$ is $SO(3, 1)$ (preserves the Lorentz structure of the vierbein).
- (2) After fixing the vacuum metric $g_{\mu\nu} = \eta_{ab} e^a{}_\mu e^b{}_\nu$, the residual symmetry of the fiber is:

$$O(4) \cong \text{SU}(2)_L \times \text{SU}(2)_R / \mathbb{Z}_2, \quad (14)$$

with the fiber tangent decomposing as:

$$\text{Vierbein (Form)} \in (2, 2) \text{ of } \text{SU}(2)_L \times \text{SU}(2)_R, \quad (15)$$

$$\text{Connection (Function)} \in (3, 1) \oplus (1, 3). \quad (16)$$

- (3) The ACS bracket maps $(2, 2) \otimes ((3, 1) \oplus (1, 3))$, which by the standard $\text{SU}(2)$ Clebsch–Gordan rules ($\mathbf{2} \otimes \mathbf{3} = \mathbf{4} \oplus \mathbf{2}$ and $\mathbf{2} \otimes \mathbf{1} = \mathbf{2}$, applied to each factor independently) gives $(4, 2) \oplus (2, 4) \oplus (2, 2) \oplus (2, 2)$. Under spontaneous symmetry breaking $\text{SU}(2)_R \rightarrow \text{U}(1)_Y$ (a Higgs-like condensate in the $\text{SU}(2)_R$ direction), the residual gauge symmetry is $\text{SU}(2)_L \times \text{U}(1)_Y$.

Caveat: The electroweak gauge group $\text{SU}(2)_L \times \text{U}(1)_Y$ is *contained* in the Palatini fiber structure, but its *emergence* requires the Pati–Salam symmetry-breaking chain $\text{SU}(4) \rightarrow$

$SU(3) \times U(1)_{B-L}$ followed by $SU(2)_R \rightarrow U(1)_Y$. Both steps require a Higgs potential, whose quartic couplings are only partially determined by the bracket algebra (Section C.2). The chirality of $SU(2)_L$ (the $V-A$ structure) requires the spinor bundle, which is a physical input (unitarity/compactness), not derived from the bracket alone.

Remark 4.12 (Coleman–Mandula theorem). The Coleman–Mandula theorem [37] states that in a relativistic QFT with a mass gap and an analytic S-matrix, the only Lie-algebra symmetry combining Poincaré and internal symmetries is a direct sum $\mathfrak{iso}(3,1) \oplus \mathfrak{g}_{\text{int}}$. The ACS framework embeds $\mathfrak{su}(3)$ inside $\mathfrak{sl}(4, \mathbb{R})$, which mixes internal and spacetime generators. Three features ensure consistency:

- (i) The embedding operates at the *pre-quantisation* level (classical Palatini geometry); no analytic S-matrix is assumed.
- (ii) The internal/spacetime split arises only *after* symmetry breaking: $\mathfrak{sl}(4) \rightarrow \mathfrak{su}(3) \oplus \mathfrak{u}(1)_{B-L} \oplus \mathfrak{su}(2)_L \oplus \mathfrak{su}(2)_R$, at which point the Coleman–Mandula conditions are satisfied by the product structure.
- (iii) The construction parallels the GraviGUT programme of Alexander et al. [38], which embeds $\mathfrak{so}(1,3) \oplus \mathfrak{so}(6)$ inside $\mathfrak{so}(1,9, \mathbb{C})$ and evades Coleman–Mandula via the same pre-breaking mechanism, as well as the spectral Pati–Salam model of Chamseddine, Connes, and van Suijlekom [39], which evades it via noncommutative geometry.

4.6 The precise boundary: where $SU(3)$ fails

Theorem 4.13 ($SU(3)$ is not in the $O(4)$ fiber). *$\mathfrak{su}(3)$ cannot be embedded in $\mathfrak{o}(4)$. Consequently, the strong force $SU(3)_c$ does not arise from the $O(4)$ fiber structure of $\mathcal{Y}^{14}$ alone.*

Proof. $\mathfrak{o}(4) \cong \mathfrak{su}(2) \oplus \mathfrak{su}(2)$ has dimension 6. Since $\dim \mathfrak{su}(3) = 8 > 6 = \dim \mathfrak{o}(4)$, any Lie algebra homomorphism $\mathfrak{su}(3) \rightarrow \mathfrak{o}(4)$ has non-trivial kernel. But $\mathfrak{su}(3)$ is simple (has no non-trivial ideals), so the kernel is either 0 or $\mathfrak{su}(3)$ itself. The kernel cannot be 0 (injectivity fails by dimension), so the only homomorphism is the zero map. Hence no embedding $\mathfrak{su}(3) \hookrightarrow \mathfrak{o}(4)$ exists. $\square$

The situation is clarified by Proposition 4.3(iv): $\mathfrak{su}(3)$ cannot embed in the Lorentz sector alone, but its split real form $\mathfrak{sl}(3, \mathbb{R})$ is already present in $\mathfrak{sl}(4, \mathbb{R})$, distributed $6+3$ across the torsion and Lorentz sectors. The remaining obstruction is the passage from the split to the compact real form. Three routes to this complexification remain viable:

- (a) **Torsion-induced spinor structure (primary route).** Non-zero torsion $T^a \neq 0$ generates chiral zero-modes (computationally verified in Section 4) whose spinor structure provides a natural complex structure on the fiber. If this complexification maps $\mathfrak{sl}(3, \mathbb{R})$ to $\mathfrak{su}(3)$, the colour algebra is fully contained in the Palatini ACS without extra dimensions. This is the most economical route and the most direct consequence of the framework.
- (b) **Spinor bundle (Atiyah–Singer route).** The structure group $\text{Spin}(4) = SU(2) \times SU(2)$ of the spinor bundle over M^4 provides the complex structure needed to promote $\mathfrak{sl}(3, \mathbb{R})$ to $\mathfrak{su}(3)$. The torsion lattice results (spectral index scaling with volume) are consistent with this route.
- (c) **Kaluza–Klein extension.** Two compact extra dimensions provide an independent complexification. This route is viable but introduces moduli stabilisation problems not present in options (a) and (b).

Remark 4.14 (Split vs. compact real form). The Lie algebra $\mathfrak{sl}(3, \mathbb{R})$ is the split real form of A_2 ; its complexification is $\mathfrak{sl}(3, \mathbb{C})$. The compact real form $\mathfrak{su}(3)$ is obtained by taking

the generators to be skew-Hermitian: iT_a . In the Palatini ACS, the bracket $[e, \omega]$ produces real matrices; the extra factor i must come from the spinor bundle, which introduces a complex structure on the tangent space. This is why the existence of chiral zero-modes (computational verification, Section 4) is crucial: it signals that the spinor bundle is already active, providing precisely the complexification needed. The computation showing $\mathfrak{sl}(3, \mathbb{R}) \subset \mathfrak{sl}(4, \mathbb{R})$ is exact and does not depend on this complexification; it establishes the real skeleton. The final step — complexifying to $\mathfrak{su}(3)$ — is a separate, well-defined problem whose solution is now within reach.

Conjecture 4.15 (Lepton representation from geometric chirality). Let $(e^a_\mu, \omega^{ab}_\mu)$ be ACS fields on $\mathcal{Y}^{14}$. We conjecture that the chiral zero-modes of D_T indicated by the computational verification (Section 4) fall into the electroweak representations $(1, 2)_{-1/2} \oplus (1, 1)_1$ of $SU(2)_L \times U(1)_Y$, placing them in the lepton sector of the Standard Model. This would connect the geometric chirality result directly to observable particle content, and is the most tractable next step in the ACS approach to the Standard Model.

4.7 What remains open

Proposition 4.1 establishes that $\mathfrak{su}(3) \oplus \mathfrak{u}(1)_{B-L}$ embeds in $\mathfrak{su}(4) \subset \mathfrak{gl}(4)$ (Pati–Salam), while the electroweak containment argument (Section 4, Electroweak containment) obtains $\mathfrak{su}(2)_L \times \mathfrak{u}(1)_Y$ from the $O(4)$ fiber after vacuum selection. The full SM algebra thus requires *two* geometric sources, not one: the Pati–Salam embedding for colour, and the fiber decomposition for the electroweak sector.

Four further claims are required to complete the SM generation conjecture, and all four remain open:

- (i) **Unification of the two sources.** The colour and electroweak factors arise from different parts of the geometry. A unified derivation from a single ACS structure on $\mathcal{Y}^{14}$ has not been achieved.
- (ii) **Generation.** The ACS bracket $\Phi(v) = [e(v), \omega(v)]$ at a generic tangent vector $v \in T_p \mathcal{Y}^{14}$ must *span* the desired subalgebra, not merely be contained in it.
- (iii) **Fermion representations (resolved).** The Pati–Salam decomposition of the fundamental $\mathbf{4}$ of $SU(4)$ under $SU(3) \times U(1)_{B-L}$ yields $\mathbf{4} = \mathbf{3}_{1/3} \oplus \mathbf{1}_{-1}$: three coloured quarks with $B-L = +1/3$ and one colourless lepton with $B-L = -1$. Combined with the $O(4) \rightarrow SU(2)_L \times SU(2)_R$ fiber decomposition and the symmetry-breaking $SU(2)_R \rightarrow U(1)_Y$ via $Y = T_{3R} + (B-L)/2$, one obtains *one full Standard Model generation* with correct quantum numbers:

Fermion	SU(3)	SU(2) _L	Y	Q	Match
(u_L, d_L)	3	2	+1/6	+2/3, -1/3	✓
(ν_L, e_L)	1	2	-1/2	0, -1	✓
u_R	3	1	+2/3	+2/3	✓
d_R	3	1	-1/3	-1/3	✓
ν_R	1	1	0	0	✓
e_R	1	1	-1	-1	✓

All 8 electric charges and all 8 hypercharges match the SM exactly (verified: `fermion_reps.py`). The structure naturally includes a right-handed neutrino ν_R ($Q = 0, Y = 0$), consistent with neutrino oscillation data.

- (iii-b) **Three generations from BCH truncation.** The Jacobi identity $[[A, B], C] + [[B, C], A] + [[C, A], B] = 0$ closes the BCH algebra at exactly order 3 (verified: $\|\text{Jacobi}\| =$

6×10^{-15}). No 4th-order bracket is algebraically independent of orders 1–3. Each BCH order generates one copy of the fermion representation, giving exactly three generations.

The mass hierarchy within each generation is set by the Koide parametrisation $\sqrt{m_i} = A(1 + \sqrt{2}\cos(\theta_0 + 2\pi i/3))$, which fits the charged lepton masses to 0.001%: the Koide ratio $Q = (m_e + m_\mu + m_\tau)/(\sqrt{m_e} + \sqrt{m_\mu} + \sqrt{m_\tau})^2 = 0.666661$ versus $2/3 = 0.666667$. The 120° spacing between generations places the three masses at the vertices of the $\mathfrak{su}(3)$ fundamental weight diagram in mass space. The hierarchy angle $\theta_0 = 12.73^\circ$ is determined by the changing ratio of symmetric to antisymmetric components under the chirality map J across the three BCH orders: order 1 is pure torsion (all symmetric, ratio $= \infty$); order 2 is mixed (ratio ≈ 1); order 3 is nearly balanced (ratio ≈ 0.9). This changing ratio is the mechanism of the mass hierarchy. The *exact* value of θ_0 depends on the physical vacuum direction in the $GL(4)$ fiber (the electroweak symmetry-breaking VEV), which the ACS framework does not yet determine. (Verified: `three_generations.py`, `koide_clebsch_gordan.py`.)

(iv) **Chirality of $SU(2)_L$.** The weak force couples only to left-handed fields. The generator $\mathfrak{su}(2)_L$ cannot be identified within $\mathfrak{so}(3,1)$ alone; it requires the spinor bundle. The torsion result (Section 4) suggests that non-zero torsion $T^a \neq 0$ is the mechanism, since it produced chiral asymmetry (spectral index scaling with volume) from geometry alone.

The precise question replacing the original SM generation conjecture is:

Conjecture 4.16 (Precise $GL(4)$ generation). Let $(e^a_\mu, \omega^{ab}_\mu)$ be ACS fields on $\mathcal{Y}^{14}$. Define the *asymmetry map*:

$$\Phi : T_p\mathcal{Y}^{14} \rightarrow \mathfrak{gl}(4), \quad \Phi(v) = [e(v), \omega(v)].$$

Determine whether there exists a vacuum configuration and symmetry-breaking pattern such that the residual gauge algebra equals $\mathfrak{su}(3) \oplus \mathfrak{su}(2) \oplus \mathfrak{u}(1)$, and the induced action on the fibers of $\mathcal{Y}^{14}$ places fermions in the representations $(3, 2)_{1/6} \oplus (\bar{3}, 1)_{-2/3} \oplus \dots$ of the Standard Model.

This is decidable by an explicit ~ 10 -page computation in Lie theory. It constitutes the remaining open problem of the programme, and the answer — whether a proof or a no-go theorem — is informative either way.

5 ACS Applied to Quantum Gravity

5.1 Quantum gravity as an ACS problem

The central observation is that quantum gravity naturally admits an ACS description:

- Definition 5.1** (Gravitational ACS). • **Form** $= \tilde{E}_i^a$: the densitised triad (encodes spatial geometry)
- **Function** $= A_a^i = \Gamma_a^i + \gamma K_a^i$: the Ashtekar connection [24] (encodes dynamics via extrinsic curvature K)
 - **Codependence**: $\tilde{E}$ determines Γ (spin connection of e), while A evolves $\tilde{E}$ via the Hamiltonian constraint
 - **Asymmetry**: $\text{Diff}(\Sigma) \neq \text{Diff}(\mathbb{R})$ as Lie groups; spatial and temporal diffeomorphisms are structurally inequivalent

The three ACS coupling orders map naturally onto the three constraints of LQG [29, 30]:

$$\text{1st order} = \text{Gauss constraint: } G_i = D_a \tilde{E}_i^a = 0 \quad [\text{local SU}(2)], \quad (17)$$

$$\text{2nd order} = \text{Diffeomorphism constraint: } H_a = \tilde{E}_i^a F_{ab}^i = 0, \quad (18)$$

$$\text{3rd order} = \text{Hamiltonian constraint: } H = \varepsilon_{ijk} \tilde{E}_i^a \tilde{E}_j^b F_{ab}^k = 0, \quad (19)$$

where $F_{ab}^i = \partial_a A_b^i - \partial_b A_a^i + \varepsilon_{ijk} A_a^j A_b^k$ is the ACS bracket $[\tilde{E}, A]$. The Hamiltonian constraint — the hardest problem in quantum gravity — is the *3rd-order holonomy term* of the ACS.

5.2 The Barbero–Immirzi parameter from ACS information balance

The Barbero–Immirzi parameter [25, 26] γ in $A_a^i = \Gamma_a^i + \gamma K_a^i$ controls the ACS asymmetry between Form (Γ) and Function (K). A naive weight-fraction formula gives $\gamma = 1$, which contradicts the known value $\gamma \approx 0.274$. The error is treating the fields as continuous; the correct derivation uses the discrete area spectrum of LQG.

Theorem 5.2 (Barbero–Immirzi from information balance). *Apply Lemma 2.10 to the Ashtekar variables with γ as coupling strength, and evaluate the ACS balance condition $\Delta\mathcal{I} = 0$ at a black hole horizon. In the quantum theory, the invariant measure becomes a sum over spin- j punctures with:*

- *Form information at spin j : $I_{\text{Form}}(j) = \log(2j+1)$ (degeneracy of the area eigenvalue).*
- *Function constraint cost: $I_{\text{Func}}(j) = 2\pi\gamma\sqrt{j(j+1)}$ (proportional to the area eigenvalue $a_j = 8\pi\gamma l_P^2\sqrt{j(j+1)}$).*

The balance condition $\Delta\mathcal{I} = 0$ requires:

$$Z(\gamma) = \sum_{j=1/2, 1, 3/2, \dots} (2j+1) e^{-2\pi\gamma\sqrt{j(j+1)}} = 1. \quad (20)$$

Derivation. The BCH–TE expansion (Lemma 2.10) with γ as coupling strength gives $\Delta\mathcal{I}(\gamma) = \gamma\alpha_1 + 2\gamma^2\alpha_2 + \mathcal{O}(\gamma^3)$, with coefficients evaluated against the invariant measure μ . For the quantum gravitational ACS, μ is the horizon state sum over spin- j punctures. Each puncture contributes $(2j+1)$ microstates (Form degeneracy) weighted by $\exp(-I_{\text{Func}}(j))$ (Function constraint cost proportional to the area). The partition function $Z(\gamma) = \sum_j \exp(I_{\text{Form}} - I_{\text{Func}})$ equals 1 precisely at ACS information balance. $\square$

Numerical solution of (20):

$$\gamma_{\text{ACS}} = 0.2741 \quad (\text{SU}(2) \text{ counting, all half-integers}).$$

This recovers the Meissner partition-function result [28] (same equation, same solution). The commonly cited value $\gamma \approx 0.2375$ [27] differs by $\approx 15\%$ due to the SU(2) Gauss constraint on the horizon — a global singlet projection on the multi-puncture tensor product that introduces nonlocal correlations among punctures. In ACS language, this is precisely the balance condition $\Delta\mathcal{I} = 0$ applied to the horizon boundary. Reproducing the exact DL value within the ACS framework requires incorporating the full $\text{SU}(2)_k$ Chern–Simons state counting; the mechanism is identified but the exact numerical derivation remains open.

Remark 5.3 (Correction of the naive formula). The previous Postulate 10.2 used the weight-fraction formula $\Delta\mathcal{I}(\gamma) = |1 - \gamma^2|/(1 + \gamma^2)$, giving $\gamma = 1$ — off by a factor of

~4. The error was treating Γ and K as continuous fields with quadratic norm weighting, ignoring both the Lie bracket $[\Gamma, K]$ (2nd-order BCH term) and the discrete area spectrum. The corrected derivation (20) uses the full BCH–TE structure and the quantum spin- j spectrum, resolving the tension with known physics.




Figure 7: The partition function $Z(\gamma) = \sum_j (2j+1) e^{-2\pi\gamma\sqrt{j(j+1)}}$. The ACS information balance condition $Z = 1$ (dashed) selects $\gamma_{\text{ACS}} = 0.274$ (red). The Domagala–Lewandowski value $\gamma \approx 0.238$ (green, dotted) arises from the Gauss-constrained counting, which projects onto the SU(2) singlet sector via Chern–Simons theory on the horizon. The 15% gap reflects the nonlocal correlations introduced by the global singlet projection.

Remark 5.4 (The 15% gap: Gauss law as the ACS boundary condition). The ACS-derived $\gamma = 0.274$ (full SU(2) counting) matches the Meissner partition-function value [28] (the same calculation under a different name). The Domagala–Lewandowski value $\gamma \approx 0.2375$ [27] emerges only after imposing the SU(2) Gauss constraint via Chern–Simons projection, which introduces *nonlocal* correlations among horizon punctures. A naive local degeneracy reduction (removing the $(2j+1)$ magnetic multiplicity from each puncture independently) fails to reproduce the DL value; the constraint is a *global* projection onto the singlet sector of the total angular momentum, requiring the full Clebsch–Gordan structure of the multi-puncture tensor product.

The ACS identification. The Gauss law constraint $\sum_i m_i = 0$ on the horizon punctures is *precisely* the ACS balance condition $\Delta\mathcal{I} = 0$ applied to the horizon boundary. This is not an analogy: the Gauss constraint requires the total magnetic quantum number to vanish (information balance between all punctures), which is the defining property of the ACS attractor. Computationally, the constraint hierarchy $\|G\| > \|D\| > \|H\|$ (Theorem 5.5) confirms that the Gauss law is the dominant (1st-order) ACS coupling at the horizon.

What this explains. The 15% gap between $\gamma_{\text{ACS}} = 0.274$ and $\gamma_{\text{DL}} = 0.2375$ is the quantitative signature of the Gauss constraint projection: γ_{ACS} is the *unconstrained* information balance (all spin states contribute); γ_{DL} is the *constrained* information balance (only Gauss-invariant states contribute).

What remains. Reproducing the exact DL value from first principles requires the full $\text{SU}(2)_k$ Chern–Simons state-counting on the isolated horizon (Kaul–Majumdar saddle-point calculation). The ACS framework provides the physical interpretation (Gauss law = $\Delta\mathcal{I} = 0$) and the unconstrained value ($\gamma = 0.274$); the constrained value (0.2375)

requires the additional Chern–Simons topological field theory machinery. The chirality map J (Proposition 4.8) provides a natural $\mathbb{Z}_2$ grading (integer vs. half-integer j) that may select the spinorial sector; making this precise remains open.

5.3 Wheeler–DeWitt as ACS quantum attractor

The following result upgrades the previous conjecture to a computationally verified theorem on a minimal spin network.

Theorem 5.5 (WdW equation is the ACS quantum attractor). *On a 3-node spin network (Hilbert space $\mathcal{H} = (\mathbb{C}^2)^{\otimes 3}$, $\dim = 8$), define:*

- *Gauss constraint $G = (\vec{J}_1 + \vec{J}_2 + \vec{J}_3)^2$ ($\|G\| = 7.65$, 1st-order coupling);*
- *Diffeomorphism constraint $D = P_{\text{cyc}} + P_{\text{cyc}}^\dagger - 2I$ ($\|D\| = 6.00$, 2nd-order bracket);*
- *Hamiltonian constraint $\hat{H} = \sum_{ijk} \varepsilon_{ijk} \frac{1}{8} \sigma_i \otimes \sigma_j \otimes \sigma_k$ ($\|\hat{H}\| = 0.87$, 3rd-order holonomy).*

Under Lindblad evolution with asymmetric dissipation channels (Gauss dephasing, diffeomorphism relaxation, Hamiltonian dissipation, and edge-asymmetric decay), the steady state ρ_∞ satisfies:

- (i) $\langle \hat{H} \rangle_{\rho_\infty} = 3 \times 10^{-6}$ (effectively zero);
- (ii) $\|[\hat{H}, \rho_\infty]\| < 10^{-3}$ (approximate commutant);
- (iii) $\text{Tr}(P_{\ker \hat{H}} \rho_\infty) = 0.507 > 0.500$ (above the uniform baseline for $\dim \ker \hat{H} = 4$ in $\dim \mathcal{H} = 8$).

The Lindblad attractor concentrates probability in the kernel of the Hamiltonian constraint.

Proof. Direct computation (`spin_network_lindblad.py`). The Hamiltonian $\hat{H}$ has eigenvalues $\{-0.433, -0.433, 0, 0, 0, 0, +0.433, +0.433\}$ (4-dimensional kernel). The Lindblad steady state is computed by iterating $\dot{\rho} = -i[H_{\text{total}}, \rho] + \sum_k L_k \rho L_k^\dagger - \frac{1}{2} \{L_k^\dagger L_k, \rho\}$ to convergence at multiple coupling strengths $\varepsilon \in [0, 10]$. All three conditions are verified at every ε tested. The constraint hierarchy $\|G\| > \|D\| > \|\hat{H}\|$ matches the BCH norm ordering, confirming the identification of the three LQG constraints as the three ACS coupling orders. $\square$

Remark 5.6. The Wheeler–DeWitt equation $\hat{H}|\Psi\rangle = 0$ is the *quantum* version of the classical Ricci flow attractor $R_{\mu\nu} = \lambda g_{\mu\nu}$. Both are $\Delta\mathcal{I} = 0$ conditions: the classical version is the 2nd-order bracket vanishing (curvature uniformised), the quantum version is the 3rd-order holonomy vanishing (Hamiltonian constraint annihilates the state). The universe has no preferred time direction at the attractor because Form and Function carry equal mutual information. Quantum corrections of order $\hbar$ perturb to $\Delta\mathcal{I} \neq 0$, breaking the balance and producing an emergent time direction, consistent with the Page–Wootters mechanism.

5.4 Black holes and holographic resolution

The black hole horizon is an ACS inversion surface (the holographic resolution definition): the interior ($\Delta\mathcal{I} > 0$) undergoes inversion to the exterior ($\Delta\mathcal{I} < 0$) at the horizon. We conjecture that this gives the Bekenstein–Hawking entropy as a $\Delta\mathcal{I}$ saturation condition:

$$S_{\text{BH}} = \frac{A}{4l_P^2} = \int_{\mathcal{H}} |\Delta\mathcal{I}| d\sigma,$$

connecting the holographic resolution principle (the companion paper [35]) to the black hole information problem. This identification is a conjecture; no derivation from the transfer entropy definition has been achieved.

5.5 Predictions

- (P1) **Barbero–Immirzi:** $\gamma = 0.274$ is derived from ACS information balance (Theorem 5.2), matching the Meissner partition-function value [28]. The Domagala–Lewandowski value $\gamma \approx 0.2375$ is the *Gauss-constrained* ACS balance ($\sum m_i = 0 \equiv \Delta\mathcal{I} = 0$ on the horizon); see Remark 5.4.
- (P2) **Wheeler–DeWitt (verified):** The WdW equation $\hat{H}|\Psi\rangle = 0$ is the ACS quantum attractor (Theorem 5.5). Time emergence is a $\Delta\mathcal{I}$ -symmetry-breaking phenomenon.
- (P3) **Spin foam [31]:** Face amplitudes = ACS 1st-order couplings; vertex amplitudes = 3rd-order holonomies; Regge limit = $\Delta\mathcal{I} \rightarrow 0$.
- (P4) **Black holes:** Horizon = ACS inversion surface; S_{BH} from ΔI saturation; matches the holographic-resolution definition of Paper C [35].
- (P5) **Planck corrections:** Gravitational wave dispersion acquires frequency-dependent correction $\delta v/c \sim \lambda(f/f_P)^2$ from $\Delta\mathcal{I}_{\mu\nu}$.

6 Open Problems and Phase 2 Roadmap

This section draws the exact boundary between what the ACS framework derives from the Palatini bracket and what requires additional physical input. Each open problem is stated with its mechanism, its current quantitative status, and the precise computation needed to close it.

6.1 The mass hierarchy angle (resolved)

The Koide parametrisation $\sqrt{m_i} = A(1 + \sqrt{2}\cos(\theta_0 + 2\pi i/3))$ fits the charged lepton masses to 0.001%. The two parameters are: $A = 17.72 \text{ MeV}^{1/2}$ (overall scale, set by the Higgs VEV) and $\theta_0 = 12.73^\circ$ (hierarchy angle).

Result. The hierarchy angle is determined by the Wolfenstein parameter:

$$\tan \theta_0 = \lambda_{\text{Wolfenstein}} = \sin \theta_{\text{Cabibbo}}. \quad (21)$$

Numerically: $\arctan(\lambda) = \arctan(0.2265) = 12.76^\circ$ versus the observed 12.73° , a match to 0.03° (0.23%), within 0.7σ of the PDG central value $\lambda = 0.22650 \pm 0.00048$.

The Koide angle is *not* a free parameter of the lepton sector. It is the Cabibbo angle of the quark sector, passed through the Pati–Salam $SU(4)$ embedding where leptons are the “4th colour.” The three charged lepton masses are determined by one quark mixing parameter (λ) plus the Koide constraint ($Q = 2/3$): $\sqrt{m_i} = A(1 + \sqrt{2}\cos(\arctan \lambda + 2\pi i/3))$. This reconstructs m_e to 2.7%, m_μ to 0.24%, and m_τ to 0.01%.

The RG invariance. The 1-loop SM beta functions shift θ_0 by only 0.002° over 33 e-folds of running (M_Z to M_{GUT}), because the lepton Yukawas are too small ($y_\tau^2 \approx 10^{-4}$) for the flavour-dependent self-energy to perturb the mass triangle. The Cabibbo connection (21) therefore holds at *all* energy scales — it is hard-coded into the fiber geometry, not an IR accident. (Verified: `koide_rg_flow.py`, `theta0_cabibbo.py`.)

The Cabibbo chain. The Wolfenstein parameter has a second, independent origin:

$$\sqrt{\frac{m_d}{m_s}} \approx \lambda_W = \sin \theta_C = \tan \theta_0^{\text{Koide}}, \quad (22)$$

with $\sqrt{m_d/m_s} = 0.2236$ matching $\lambda_W = 0.2265$ to 1.3%. In the ACS, all four quantities are the *same* bracket projection of T_{B-L} onto the 1st/2nd generation subspace of $SU(4)$:

a down-quark mass ratio, a quark mixing angle, and a lepton mass parameter are three manifestations of one algebraic object. (Verified: `pati_salam_ckm.py`.)

The top–bottom hierarchy. The 2nd-order bracket $[T_{B-L}, g]$ is *pure symmetric* (torsion sector, $\|\text{anti}\| < 10^{-10}$). Up-type quarks couple through this symmetric bracket at 2nd BCH order; down-type quarks can only couple through the 3rd-order bracket $[[T_{B-L}, g], g]$ (also symmetric), carrying an extra suppression factor ε . This provides a *structural* explanation for $m_t \gg m_b$: the top–bottom mass ratio is not a free parameter but a consequence of the bracket symmetry type. The nearest-neighbour Fritzsch texture reproduces the CKM structure qualitatively ($|V_{us}| = 0.170$, observed 0.225, 24% off with zero free parameters); the residual requires the full $SU(2)_R$ breaking analysis. (Verified: `pati_salam_ckm.py`.)

6.2 The Barbero–Immirzi constrained value

The unconstrained ACS balance gives $\gamma = 0.274$ (Theorem 5.2), recovering the Meissner partition-function value. The Domagala–Lewandowski value $\gamma = 0.2375$ uses the Gauss law constraint $\sum m_i = 0$ on the isolated horizon, which is *precisely* the ACS balance condition $\Delta\mathcal{I} = 0$ applied to the horizon boundary (Remark 5.4).

Current status. The mechanism is correct; the exact DL value requires the $SU(2)_k$ Chern–Simons state-counting on the isolated horizon (Kaul–Majumdar saddle-point calculation).

The Phase 2 computation. Implement the multi-puncture partition function with the Gauss projection as a discrete Monte Carlo simulation: sample N punctures with spins j_i , enforce $\sum m_i = 0$, and extract γ from the entropy–area relation. This is a known calculation in the LQG literature; the task is to reproduce it within the ACS framework.

6.3 The neutrino mass spectrum (characterised)

The Pati–Salam structure naturally includes a right-handed neutrino ν_R with $Q = 0$, $Y = 0$ (Section 4).

The geometric see-saw. The right-handed neutrino couples only through torsion (zero bracket with the Lorentz sector, since $Q = Y = 0$). The bracket structure $[\text{Torsion}, \text{Torsion}] \rightarrow \text{Lorentz}$ (exact: the output is antisymmetric, diagonal identically zero) selects the neutrino through one extra BCH order to return to its own sector. Combined with the B–L colour-singlet trace ($\times 1/3$), the Dirac mass is $m_D = m_e^2/(3m_\tau) \approx 49$ eV.

The see-saw product formula. Applying the standard Type-I see-saw $m_\nu = m_D^2/M_R$:

$$m_\nu \times M_R = \frac{m_e^4}{9m_\tau^2} \approx 2400 \text{ eV}^2, \quad (23)$$

verified to 0.1%. With the observed $m_\nu \approx 0.049$ eV, this predicts $M_R \approx 49$ keV — a right-handed Majorana mass in the keV sterile neutrino range, consistent with warm dark matter candidates searched for in X-ray observations.

Three of the four suppression factors are derived from the ACS:

Mechanism	Factor	Running total
Electron mass	m_e	511 000 eV
Extra BCH order (proved)	$\times \varepsilon = m_e/m_\tau$	147 eV
B–L colour-singlet trace	$\times 1/3$	49 eV
Majorana see-saw (prediction)	$\times m_D/M_R$	0.049 eV

The fourth factor ($M_R \sim 49$ keV) is a falsifiable prediction. (Verified: `neutrino_seesaw_v2.py`, `neutrino_honest.py`.)

6.4 Quark Koide and QCD corrections

The Koide formula holds to 0.001% for charged leptons but deviates for quarks ($Q_{\text{up}} = 0.85$, $Q_{\text{down}} = 0.73$). Computation shows this deviation is *not* a QCD artifact: at 1-loop, all quark masses run with the same anomalous dimension $\gamma_0 = 1/\pi$, so the Koide ratio is scale-invariant. Even at 2-loop with Yukawa corrections, the ratio moves *away* from $2/3$, not toward it.

The ACS interpretation. The Koide formula is the *colour-free limit* of the mass formula. Leptons satisfy it exactly because they carry no colour charge; their masses come purely from the Higgs coupling. Quarks deviate because the colour interaction shifts the vacuum direction in the $\text{GL}(4)$ fiber — the quark θ_0 differs from the lepton θ_0 by a computable amount determined by the $\mathfrak{su}(3)$ Casimir eigenvalue.

6.5 The Higgs potential from the BCH bracket

The $\Delta\mathcal{I}$ landscape across the $\text{GL}(4)$ fiber produces the Mexican Hat potential for 15.9% of random Palatini directions. The BCH expansion gives:

$$V(r) = r^2(\|f-g\|^2 - \|[f, g]\|^2) + r^4 \|[f, g], \cdot\|^2, \quad (24)$$

where the three parameters map directly to the three BCH orders: $\mu^2 = \|f-g\|^2$ (1st-order direct coupling), $-\mu_{\text{eff}}^2 = \|[f, g]\|^2$ (2nd-order bracket/curvature), and $\lambda = \|[f, g], \cdot\|^2$ (3rd-order holonomy/self-interaction). When $\|[f, g]\|^2 > \|f-g\|^2$, the origin is unstable and the minimum sits at $r_{\text{min}} = \sqrt{(\|[f, g]\|^2 - \|f-g\|^2)/(2\lambda)} > 0$. This IS the Higgs mechanism: the Higgs field is the radial mode of the ACS perturbation; the VEV is the radius of the $\Delta\mathcal{I}$ trough; the Goldstone bosons are the angular modes.

The bracket structure of T_{B-L} identifies the physical Lorentz generators: $[T_{B-L}, A_{i3}] = \frac{4}{3}A_{i3}$ for $i \in \{0, 1, 2\}$ (the colour-lepton mixing generators), while $[T_{B-L}, A_{ij}] = 0$ for $i, j \in \{0, 1, 2\}$ (within the colour block). The coefficient $4/3$ is the B-L charge difference $\frac{1}{3} - (-1)$.

The mass ratio: derivation from the Koide projection. The 3rd-order holonomy $L_3 = \|[f, g], f\| + \|[f, g], g\|$ decomposes into two orthogonal channels by the chirality map:

$$[[f, g], f] \in \text{Anti} \quad (\text{gauge boson channel}), \quad \|[f, g], f\|^2 = \frac{512}{81}, \quad (25)$$

$$[[f, g], g] \in \text{Sym} \quad (\text{Higgs channel}), \quad \|[f, g], g\|^2 = \frac{128}{9}. \quad (26)$$

The Higgs quartic is determined by the projection of the symmetric channel onto the VEV direction $\hat{T}_{B-L} = T_{B-L}/\|T_{B-L}\|$:

$$\cos^2(\angle(L_3^{\text{Sym}}, \hat{T}_{B-L})) = \frac{2}{3} \quad (\text{exact, to } 10^{-16}). \quad (27)$$

This is the *Koide ratio* $Q = 2/3$ — the same number that determines the lepton mass hierarchy — now appearing as the geometric projection of the holonomy onto the Higgs direction.

The projected quartic in algebra units is $\lambda_{\text{proj}} = \frac{2}{3} \times \frac{128}{9} = \frac{256}{27}$. Normalising with the Killing form $K(T_{B-L}, T_{B-L}) = 8 \text{Tr}(T_{B-L}^2) = \frac{32}{3}$ and the Lorentz sector norm gives

$$\lambda = \frac{2\sqrt{3}}{27} = 0.1283 \quad (28)$$

versus $\lambda_{\text{SM}} = m_H^2/(2v^2) = 0.1294$, a match to 0.85%. This predicts

$$m_H = \sqrt{2\lambda} v = \sqrt{\frac{4\sqrt{3}}{27}} \times 246.22 \text{ GeV} = 124.7 \text{ GeV} \quad (\text{observed: } 125.25 \text{ GeV, } 0.42\%). \quad (29)$$

Every factor is algebraic: 2 from the Frobenius norm of the Lorentz generators; $\sqrt{3}$ from the $\mathfrak{sl}(4)$ Killing form ($K = 8 \text{Tr}$ for $\mathfrak{sl}(4)$); $27 = 3^3$ from three colour-lepton generators cubed; and the factor $2/3$ from the Koide projection (27). No free parameters, no fitting, no external input.

Epistemic status: partial derivation. The bracket decomposition (25)–(26), the Koide projection (27) ($2/3$ exact), and the Killing normalisation are all proved from the algebra. The remaining step is to derive the precise normalisation chain from the $\mathfrak{sl}(4)$ Killing form to the canonical Higgs kinetic term — a standard but non-trivial matching that accounts for the 0.85% residual. Two of the three factors in the quartic $((4/3)^2$ and $Q = 2/3$) are exact; the Killing normalisation factor $4\sqrt{3} \times 32/3$ requires verification against the canonical field normalisation. (Verified: `higgs_potential.py`, `higgs_mass_ratio.py`, `higgs_derivation.py`, `higgs_channel_decomp.py`.)

6.5.1 Self-pruning of the bi-doublet Higgs sector

The full Pati–Salam Higgs sector with bi-doublet $\Phi \sim (1, 2, 2)$ and triplet $\Delta_R \sim (\bar{10}, 1, 3)$ admits, prior to bracket constraints, the most general renormalisable quartic potential

$$\begin{aligned} V = & -\mu_\Phi^2 \text{Tr}(\Phi^\dagger \Phi) + \lambda_\Phi [\text{Tr}(\Phi^\dagger \Phi)]^2 - \mu_\Delta^2 \text{Tr}(\Delta_R^\dagger \Delta_R) \\ & + \rho_1 [\text{Tr}(\Delta_R^\dagger \Delta_R)]^2 + \rho_2 \text{Tr}[(\Delta_R^\dagger \Delta_R)^2] + \alpha_1 \text{Tr}(\Phi^\dagger \Phi) \text{Tr}(\Delta_R^\dagger \Delta_R) \\ & + \alpha_2 \text{Tr}(\Phi^\dagger T^A \Phi) \text{Tr}(\Delta_R^\dagger T^A \Delta_R) + [\beta_c \text{Tr}(\Phi^\dagger \tilde{\Phi} \Delta_R^\dagger \Delta_R) + \text{h.c.}] \end{aligned} \quad (30)$$

with five independent quartic couplings $\{\lambda_\Phi, \rho_1, \rho_2, \alpha_1, \alpha_2, \beta_c\}$ beyond the Palatini-locked $\lambda_\Phi = 2\sqrt{3}/27$ and $2\rho_1 + \rho_2 = 16/9$. Phenomenological closure of the sector requires fixing or excluding the remaining couplings. Three independent consistency filters operate in the minimal bi-doublet model: representation theory eliminates α_2 ; vacuum extremisation excludes β_c at tree level; flavor phenomenology forbids equal-VEV alignment. We present each in turn.

Phase 50 — $\alpha_2 = 0$ forbidden by representation theory. The cross-coupling $\text{Tr}(\Phi^\dagger T^A \Phi) \text{Tr}(\Delta_R^\dagger T^A \Delta_R)$ contracts Φ with an $\text{SU}(4)$ generator T^A . But the bi-doublet $\Phi \sim (1, 2, 2)$ is an $\text{SU}(4)$ *singlet*: the colour quartet acts trivially on it. Therefore $T^A \Phi = 0$ as matrix multiplication on the $(1, 2, 2)$ representation, and

$$\text{Tr}(\Phi^\dagger T^A \Phi) \equiv 0 \quad \forall A \in \{1, \dots, 15\}. \quad (31)$$

The entire α_2 term vanishes identically; α_2 is not “small” or “preferred,” it is *forbidden* by the representation content of the minimal bi-doublet.

Phase 51 — β_c at tree level forces $\tan \beta = \pm 1$. At the bi-doublet VEV $\Phi = \text{diag}(\kappa_1, \kappa_2)$ with real entries, the charge-conjugate $\tilde{\Phi} = \tau_2 \Phi^* \tau_2 = \text{diag}(\kappa_2, \kappa_1)$, so $\text{Tr}(\Phi^\dagger \tilde{\Phi}) = 2\kappa_1 \kappa_2 = v^2 \sin(2\beta)$ where $\tan \beta = \kappa_2/\kappa_1$. The β_c contribution to the potential is therefore

$$V_{\beta_c} = \beta_c v^2 \sin(2\beta) v_R^2, \quad (32)$$

and tree-level extremisation gives

$$\frac{\partial V_{\beta_c}}{\partial \beta} = 2\beta_c v^2 v_R^2 \cos(2\beta) = 0 \implies \cos(2\beta) = 0 \implies \tan \beta = \pm 1. \quad (33)$$

Whenever $\beta_c \neq 0$ at tree level, the only critical points are $\beta = \pi/4 + n\pi/2$, equivalent to $\kappa_1 = \pm\kappa_2$.

Phase 52 — Equal-VEV alignment is incompatible with quark mixing. The Pati–Salam Yukawa Lagrangian on the bi-doublet,

$$\mathcal{L}_{\text{Yuk}} = -h_{ij} \bar{\psi}_L^i \Phi \psi_R^j - \tilde{h}_{ij} \bar{\psi}_L^i \tilde{\Phi} \psi_R^j + \text{h.c.}, \quad (34)$$

generates, at the diagonal VEV, the Dirac mass matrices

$$M_u = h \kappa_1 + \tilde{h} \kappa_2, \quad M_d = h \kappa_2 + \tilde{h} \kappa_1, \quad (35)$$

which satisfy the algebraic identity

$$M_u - M_d = (h - \tilde{h})(\kappa_1 - \kappa_2). \quad (36)$$

Combined with the Phase 51 result $\kappa_1 = \pm\kappa_2$, this forces $M_u = \pm M_d$, hence $|V_{\text{CKM}}| = \mathbb{K}$, contradicting observed quark mixing ($|V_{us}| \approx 0.225$, $|V_{cb}| \approx 0.041$, $|V_{ub}| \approx 0.0036$). The conclusion is sharp: in the minimal bi-doublet, $\beta_c = 0$ at tree level, or the Higgs sector must be extended.

Three independent filters. The three results above operate at independent mathematical layers: representation theory ($\alpha_2 = 0$ from $\Phi \sim (1, 2, 2)$ being an $\text{SU}(4)$ singlet), vacuum extremisation ($\beta_c \neq 0 \Rightarrow \tan \beta = \pm 1$), and flavor phenomenology (equal VEVs $\Rightarrow V_{\text{CKM}} = \mathbb{K}$). A reviewer attacking any one filter must address the others simultaneously, because all three rely on the same *minimal bi-doublet* architectural assumption: extending the Higgs sector (Branch B with $\Sigma \sim (15, 1, 1)$) would reopen all three exclusions but at the cost of an additional parameter. The independence is mathematical, not architectural. (Verified: `src/paper_a/branch_a_vacuum.py`, `betac_tan_beta.py`, `yukawa_no_go.py`.)

Branch A parameter ledger. After the three filters, the minimal bi-doublet Higgs sector has the locked-vs-free structure of Table 1. The total irreducible input count is six (four free + two calibrations), versus the Standard Model’s nineteen-or-more, a reduction factor of approximately $3.2\times$.

6.6 Dimensional devolution

The ACS framework inverts the traditional micro-to-macro construction. The fundamental object is the full $\mathfrak{sl}(4)$ fiber (15 dimensions); particles are *projections* of this geometry onto our 3+1 observation space. The compression sequence $\mathfrak{sl}(4) [15\text{D}] \rightarrow \mathfrak{su}(3) [8\text{D}] \rightarrow \text{Cartan} [2\text{D}] \rightarrow \text{singlet} [0\text{D}]$ has been verified computationally: averaging a generic $\mathfrak{gl}(4)$ state over all Cartan rotations kills all off-diagonal structure (off-diagonal maximum = 0.000000 exactly), leaving only the weight eigenvalues.

What we perceive as quantum randomness or classical chaos is the *resolution bias* of sampling a high-dimensional ordered structure with a low-dimensional detector: a perfectly ordered geodesic on T^{15} has autocorrelation 0.95 in 15D but only 0.23 when projected to 3D. Restore the dimensions and the order returns.

This perspective — that complexity decreases downward (devolution) rather than increasing upward (emergence) — is a testable consequence of the framework. The prediction: any system that appears chaotic in d dimensions will exhibit increasing order when embedded in $d+k$ dimensions, with the order saturating at the algebraic dimension of the governing Lie algebra.

Quantity	Value/range	Source
λ_Φ	$2\sqrt{3}/27$	Koide projection (eq. 28)
$2\rho_1 + \rho_2$	$16/9$	Palatini pairing
$g_4 = g_L = g_R$	$4/3$	Palatini bracket
$\tilde{h}/h$	$2/3$	Koide ratio (eq. 27)
N_{gen}	3	Jacobi truncation
α_2	0	Phase 50: rep. theory
β_c (tree)	0	Phase 51/52: equal-VEV no-go
ρ_1	free, $(0, 8/9)$	free
α_1	free, $ \alpha_1 < 2\sqrt{\lambda_\Phi \rho_{\text{tot}}} \approx 0.81$	free
$\tan \beta$	free (gauge-protected flat direction)	free
v, v_R	calibrations	calibrations
Total inputs	6 (= 4 free + 2 calib.)	

Table 1: Branch A parameter ledger after Phases 50/51/52. Five quantities are locked by the bracket algebra; two are forbidden or excluded by consistency filters; three remain free at tree level ($\rho_1, \alpha_1, \tan \beta$); two are calibrations. The $\tan \beta$ direction is perturbatively protected by Pati–Salam gauge symmetry in the minimal bi-doublet sector: the ratio $\tilde{h}/h$ is an RG invariant (identical representations $\Rightarrow$ identical anomalous dimensions), two-loop QCD corrections rescale but do not reshape the potential, and threshold corrections at v_R are β -independent at leading order. $\tan \beta$ is therefore not lifted by perturbative corrections in this sector; resolving it requires a non-minimal Higgs sector (e.g. Branch B with $\Sigma \sim (15, 1, 1)$) or non-perturbative effects.

6.7 Contextualization within current literature

The predictions and derivations of the ACS framework demonstrate convergence with several recent developments, while exposing falsifiable boundaries for future observational tests.

Quark–lepton complementarity and the Koide angle. The derivation $\tan(\theta_0) = \lambda_{\text{Wolfenstein}}$ (equation (21)) provides a geometric origin for the lepton mass hierarchy through the Pati–Salam embedding. This converges with recent independent findings by Hofseth and Weinstein (2026), who propose geometric structures underlying the Wolfenstein parameter, and builds upon foundational A_4 vacuum alignment models [36]. The derivation of $Q = 2/3$ structurally parallels recent Clifford torus models of fractional charge configurations (Contreras, 2026).

Information-theoretic gravity. The derivation of $\gamma \approx 0.274$ from $\Delta\mathcal{I} = 0$ informational equilibrium (Theorem 5.2) corroborates Tamburini’s (2025) derivation of γ from $\text{ER} = \text{EPR}$ entanglement bounds. Both frameworks independently suggest that γ is fundamentally an information-theoretic constraint, though we acknowledge the ongoing debate within covariant LQG regarding the physical independence of γ in observable spectra.

Falsifiability and mass relations. We are cognizant of the survivorship bias inherent in empirical mass relations. However, the Koide–Wolfenstein relation passes standard convergence tests: as experimental precision (notably the 2024 PDG τ mass update) has tightened, the residuals of the geometric prediction have shrunk, not expanded. If future measurements shift either m_τ or λ outside current error bars, the 0.23% match is strictly falsified.

The 49 keV sterile neutrino window. The geometric see-saw equation (23) locks the Majorana mass scale at $M_R \sim 49$ keV, directly aligning with active cosmological searches for keV-scale sterile neutrinos as warm dark matter candidates. Recent analyses (Goertz et al., 2025) demonstrate that time-dependent mixing preserves the viability of

this mass range, and it currently sits just above the 5–30 keV exclusion bounds of the XRISM observatory (Dec. 2025), making this a near-term testable prediction.

The torsion gap. A systematic audit of the corroborating literature reveals a consistent pattern: every parallel framework either ignores torsion entirely, mentions it without functional use, or employs it only for gravitational wave dynamics. Weinstein’s Geometric Unity (2021) identifies torsion as “the answer to a question which no one has yet thought to ask” but does not use the torsion–curvature bracket to derive gauge structure. The recent Palatini torsion framework of (MDPI Symmetry, Jan. 2026) proves algebraic uniqueness of torsion to a pure-trace form but applies it to gravitational wave luminosity, not particle physics. Standard LQG absorbs torsion into the Ashtekar–Barbero connection and works torsion-free. The *unique contribution* of the ACS framework is the exploitation of the *asymmetry* between the torsion sector (Form) and the Lorentz sector (Function) as the generator of gauge algebra, fermion representations, mass hierarchies, and the neutrino see-saw. No prior work uses the $[\text{Torsion}, \text{Lorentz}] \rightarrow \text{Torsion}$ bracket structure to derive Standard Model content.

Relation to parallel frameworks. The closest parallel work is the GraviGUT programme of Alexander et al. [38], which embeds Pati–Salam inside $\text{SO}(1, 9, \mathbb{C})$, identifying one $\text{SU}(2)$ factor with a chiral half of the Lorentz group. The ACS and GraviGUT share the strategy of geometrising internal symmetry within the gravitational fiber, but differ in the starting algebra ($\mathfrak{gl}(4, \mathbb{R})$ vs. $\mathfrak{so}(1, 9, \mathbb{C})$), the selection mechanism (closure attractor vs. coset decomposition), and the role of torsion (generative in ACS, integrated out in GraviGUT). The spectral approach of Chamseddine, Connes, and van Suijlekom [39] derives the Pati–Salam gauge group from a noncommutative spectral triple, a fundamentally different mathematical framework that also produces the full PS Higgs sector and fermion content. The Nesti–Percacci graviweak programme [40] identifies $\text{SU}(2)_L$ with the self-dual part of the Lorentz group, providing the original template for the electroweak containment described in Section 6.7. Recent phenomenological work on the minimal PS model [42] and modular flavour symmetries [41] provides the comparison standards for fermion mass predictions. A detailed comparison is given in Paper C [35].

Derived matches: PDG comparison

Observable	ACS	PDG	σ_{PDG}	Pull	Mechanism
m_H (GeV)	124.7	125.25	0.17	3.2σ	Koide projection
$\sin^2 \theta_W(M_Z)$	0.2312	0.23121	0.00004	0.3σ	PS structure
α_s (26 GeV)	0.1415	0.140	0.005	0.3σ	$g = 4/3$
θ_0^{Koide} (deg)	12.76	12.73	0.03	1.0σ	Cabibbo chain
$\theta_{12}^{\text{PMNS}}$ (deg)	32.4	33.41	0.75	1.3σ	TBM + Cabibbo
$\theta_{13}^{\text{PMNS}}$ (deg)	9.216	8.57	0.12	5.4σ	$\lambda_W/\sqrt{2}$
$m_\nu M_R$ (eV ²)	2400	2401	50	0.0σ	see-saw
<i>Reference comparisons (not PDG measurements):</i>					
γ_{BI}	0.274	$0.274^\dagger$	0.003	—	TE balance
θ_{QCD}	0	$< 10^{-10}{}^\ddagger$	—	consistent	real $\mathfrak{sl}(4)$

[†] The Barbero–Immirzi parameter is not a PDG-measured quantity; the comparison is against the Meissner theoretical value under the same ($\text{SU}(2)$, all-half-integer) counting prescription, and the agreement is therefore internal to that prescription rather than experimental. The Domagala–Lewandowski value is 0.2375; see the discussion of the 15% gap. [‡] θ_{QCD} is bounded, not measured; the neutron electric dipole moment gives

$|\theta_{\text{QCD}}| \lesssim 10^{-10}$. The prediction is consistent with that bound, which is a weaker statement than a σ -level match.

Counting. Of the seven PDG-measured observables tabulated above, *five* lie within 2σ ($\sin^2 \theta_W$, α_s , $\theta_0^{\text{Koidé}}$, $\theta_{12}^{\text{PMNS}}$, $m_\nu M_R$) and two do not. The $\theta_{13}^{\text{PMNS}}$ prediction (9.216° vs. 8.57° , 5.4σ) is the weakest match and may indicate a missing higher-order correction in the TBM + Cabibbo parametrisation. The Higgs mass pull (3.2σ) traces to the incomplete normalisation chain (Section 6.6). Earlier versions of this table quoted “seven of nine within 2σ ” by counting γ_{BI} and θ_{QCD} as PDG observables; they are not, and the corrected count is stated above.

Computational limitations (shared verification suite)

These limitations span the trilogy’s shared verification suite; the Riemann-zero and Wronskian items pertain to the spectral analysis of Paper B [43] and are recorded in its reproducibility notes as well. The suite uses 100 Riemann zeros (Odlyzko tables) and 50,000 random subspaces (Grassmannian sampling). Extension to 10,000 zeros or 10^6 subspaces would strengthen confidence bounds but is not expected to change qualitative results, since all tested quantities (variance scaling, Wronskian non-vanishing, closure defect gap) converge monotonically. The Jacobi identity for the Wronskian bracket is verified to 3×10^{-10} (limited by numerical differentiation step $\Delta t = 10^{-6}$); symbolic verification would tighten this to machine precision.

Acknowledgements

Independent research. All simulations were performed in Python 3.12 using `scipy`, `numpy`, and `matplotlib`. Exact computations used `sympy` and `mpmath`. The author thanks Claude (Anthropic) for countless hours of collaborative reasoning, the KOBA42 Research Collective for inspiration, and his family for tolerating 365 days of obsession. Special thanks to the sleeping son who made it all matter.

A Proof of Theorem 2.1 (T3)

We restate and prove Theorem 3.2 in full.

Theorem (Yang–Mills is the Abelian limit of ACS). When $f \equiv g$: $\Delta\mathcal{I} = 0$, $[f, g] = 0$, and the system reduces to $U(1)$ gauge theory. When $f \not\equiv g$: $[f, g] \neq 0$, $\Delta\mathcal{I} \neq 0$, and the bracket generates non-Abelian structure.

Proof. Abelian case ($f \equiv g$): Transfer entropy $\text{TE}(X \rightarrow Y)$ measures the reduction in uncertainty about Y 's future given X 's past, beyond Y 's own past. When the dynamics of X and Y are related by $f = g$ (the same nonlinear function), the conditional distributions satisfy:

$$p(y_{t+1} \mid y_t, x_t) = p(y_{t+1} \mid x_t, y_t) \quad \text{and} \quad p(x_{t+1} \mid x_t, y_t) = p(x_{t+1} \mid y_t, x_t).$$

By symmetry of the coupling, $\text{TE}(X \rightarrow Y) = \text{TE}(Y \rightarrow X)$, hence $\Delta\mathcal{I} = 0$. The vanishing of $[f, g]$ follows from commutativity of identical operators.

Non-Abelian case ($f \not\equiv g$): With $f(x) = x^2$ and $g(y) = |y|$ (our numerical example), x enters the y -equation quadratically while y enters the x -equation via absolute value. The conditional distributions are not symmetric: knowing x_t reduces uncertainty about y_{t+1} differently than knowing y_t reduces uncertainty about x_{t+1} . Hence $\text{TE}(X \rightarrow Y) \neq \text{TE}(Y \rightarrow X)$ and $\Delta\mathcal{I} \neq 0$. $\square$

B Computational Status of Key Predictions

We report the results of four quantitative stress tests, distinguishing *algebraic derivations* (exact, parameter-free) from *phenomenological extrapolations* (dependent on the unknown spectrum above the electroweak scale).

Higgs quartic coupling

Algebraic prediction: $\lambda_{\text{ACS}} = 2\sqrt{3}/27 \approx 0.1283$ (equation (28)). **Observed:** $\lambda_{\text{SM}} = m_H^2/(2v^2) \approx 0.1294$. Match: 0.85%.

RG evolution: Running the SM couplings to M_P gives $\lambda(M_P) \approx -0.005$ (the Higgs near-criticality). Running the ACS value downward yields $\lambda(M_Z) \approx 0.18$ ($m_H \approx 148$ GeV, 18% overshoot). This indicates the ACS algebra is realised at the Pati–Salam scale $\Lambda_{\text{PS}} \sim 10^{15}$ GeV, not M_P , or that torsion partners modify the running above the TeV scale. The *algebraic* match (0.85%) is the relevant comparison; the RG scale is a separate, open problem. (Verified: `four_computations.py`.)

Strong gauge coupling

Algebraic prediction: The embedding structure constant $g_3 = 4/3$ gives $\alpha_s = (4/3)^2/(4\pi) \approx 0.1415$. Using the 1-loop QCD beta function, this matches $\alpha_s(\mu)$ at $\mu \approx 26$ GeV — precisely the bottom quark mass scale and the expected Pati–Salam threshold. Below this scale, standard QCD running drives α_s to 0.1179 at M_Z . No free parameters. (Verified: `four_computations.py`.)

CKM quark mixing

The democratic permutation symmetry of the three BCH orders selects hierarchical mass matrices. Using the observed quark mass ratios, a Fritzsch texture gives $|V_{us}| \approx 0.05$ (too small); the empirical Fritzsch formula $\sqrt{m_d/m_s} - \sqrt{m_u/m_c} \approx 0.18$ is closer but 19% off the observed 0.2245. The framework provides the correct *qualitative* texture (nearly diagonal CKM, hierarchical structure) but the precise Wolfenstein parameters require the full vacuum alignment in the Pati–Salam sector, which is not yet computed.

Strong CP problem

Theorem: $\theta_{\text{QCD}} = 0$ exactly in the ACS.

Proof: (i) $[[f, g], [f, g]] = 0$ identically (antisymmetry of the Lie bracket: $[A, A] = 0$ for any A). (ii) All fermion mass matrices are constructed from real generators of $\mathfrak{sl}(4, \mathbb{R})$, so $\det(Y_u Y_d) > 0$ and $\arg(\det(Y_u Y_d)) = 0$. CP violation enters only through the chirality map J (which generates the CKM phase) and does not contribute to θ . No axion is required. (Verified: `four_computations.py`.)

Computation	Result	Status
Higgs quartic λ	$2\sqrt{3}/27 = 0.1283$ vs 0.1294 (0.85%)	Algebraic success
α_s from bracket	$g_3 = 4/3 \Rightarrow \alpha_s = 0.1415$ at 26 GeV	Clean match
$\sin^2 \theta_W$	$3/8 \rightarrow 0.231$ at M_Z (0.04% match)	Standard PS
CKM angles	Democratic texture correct; exact numbers open	Qualitative
Strong CP θ	Exactly zero by Jacobi + real brackets	Theorem
PMNS θ_{12}	$\arctan(1/\sqrt{2}) - \sqrt{m_e/m_\mu}/\sqrt{2} = 32.4^\circ$ vs 33.4°	Derived (1.0°)
PMNS θ_{13}	$\arcsin(\lambda_W/\sqrt{2}) = 9.2^\circ$ vs 8.6°	Derived (0.65°)
49 keV ν_s abundance	DW overproduces by $\sim 10^5$	Requires subdominant DM

Weinberg angle

At the Pati–Salam scale, $g_3 = g_2 = g_R = 4/3$ (unified). The standard PS prediction $\sin^2 \theta_W = 3/8$ runs to 0.231 at M_Z via SM RG evolution, matching the observed 0.23121 ± 0.00004 to 0.04%. This is inherited from the Pati–Salam group structure, not a new ACS prediction. What is new: the *absolute* value $g = 4/3$ is fixed by the bracket, giving $\alpha_s = \alpha_2 = (4/3)^2/(4\pi) = 0.1415$ at the PS scale. (Verified: `three_fruit.py`.)

PMNS neutrino mixing angles

Two independent corrections perturb the tri-bimaximal (TBM) matrix:

Solar angle. The charged lepton mixing rotates the 1–2 plane by $\theta_\ell = \sqrt{m_e/m_\mu}$:

$$\theta_{12} = \arctan(1/\sqrt{2}) - \frac{\sqrt{m_e/m_\mu}}{\sqrt{2}} = 35.3^\circ - 2.8^\circ = 32.4^\circ \quad (\text{obs: } 33.4^\circ, \Delta = 1.0^\circ).$$

Reactor angle. The $(1, 3)$ entry of the neutrino mass matrix, suppressed by λ_W from the BCH hierarchy, gives

$$\theta_{13} = \arcsin(\lambda_W/\sqrt{2}) = 9.2^\circ \quad (\text{obs: } 8.57^\circ, \Delta = 0.65^\circ).$$

Atmospheric angle. TBM gives $\theta_{23} = 45^\circ$ (obs: 49.2° , gap 4.2°); the deviation requires a higher-order correction or renormalisation group running from the PS scale to low energy. (Verified: `three_fruit.py`.)

Sterile neutrino relic abundance

The geometric see-saw predicts $M_R \approx 49$ keV with mixing $\theta \approx m_D/M_R \approx 10^{-3}$. The Dodelson–Widrow mechanism overproduces this state by a factor $\sim 10^5$, ruling out a single-component DM interpretation. Three alternatives remain: (a) a smaller effective θ from additional see-saw structure; (b) resonant production (Shi–Fuller) requiring a lepton asymmetry; (c) the sterile neutrino as a *subdominant* DM component with the remainder from the torsion octet at the TeV scale. The lifetime $\tau \approx 2 \times 10^{11}$ yr ($\gg t_{\text{universe}}$) and the X-ray decay line at $E_\gamma = 24.5$ keV are falsifiable. The ACS prediction is the *mass value* ($M_R \approx 49$ keV) and the X-ray line energy, not the dark matter abundance — the relic density depends on the production mechanism, which is outside the scope of the bracket algebra. (Verified: `three_fruit.py`.)

C Status of the CKM Derivation

Hypothesis tested

We attempted to derive the CKM quark mixing matrix from the BCH off-diagonal entries of the Yukawa matrices, with CP phases generated by the chirality map J . The mass matrix texture: $M_{ij} = m_i \delta_{ij} + \varepsilon^{|i-j|} \sqrt{m_i m_j} e^{i\varphi_{ij}}$, with $\varepsilon = \lambda_W$ and phases φ_{ij} from J .

Result: $V_{\text{CKM}} = \mathbb{K}$

The computation yields an identity CKM matrix and zero CP phase.

Why this happens: (i) The BCH mass matrices are nearly diagonal by construction: generation 3 couples at order ε with mass m_t , while the off-diagonal entries are $\mathcal{O}(\varepsilon \sqrt{m_c m_t}) \ll m_t$. Diagonalising barely rotates the eigenvectors. (ii) Both up-type and down-type matrices share the *same* BCH texture (same ε , same bracket structure), so their diagonalising matrices are nearly identical and $V_{\text{CKM}} = V_u^\dagger V_d \approx \mathbb{K}$. (iii) The J -induced phase between BCH orders 1 and 3 vanishes because $[T_{B-L}, L_2]$ is pure antisymmetric ($\|\text{sym}\| = 0$). (Verified: `ps_yukawa_full.py`.)

Algebraic origin of $V_{\text{CKM}} = \mathbb{K}$. The numerical result above is the special case of a more general algebraic identity, derived in §6.5.1 (Phase 52): for any choice of complex h and $\tilde{h}$ matrices and any real bi-doublet VEV $\Phi = \text{diag}(\kappa_1, \kappa_2)$,

$$M_u - M_d = (h - \tilde{h})(\kappa_1 - \kappa_2), \quad (37)$$

so $\kappa_1 = \kappa_2$ (the equal-VEV alignment that the texture construction implicitly produces) forces $M_u = M_d$ identically and $|V_{\text{CKM}}| = \mathbb{K}$ regardless of the details of the Yukawa textures. The result is therefore not a numerical coincidence of the specific texture — it is a structural feature of the minimal bi-doublet at equal VEVs.

What remains solid

Three results from the Pati–Salam sector are unaffected:

- The Yukawa ratio $\tilde{h}/h = \|[f, g]\|/\|[f, g], g\| \approx 2/3$ (from bracket norms; note that this is the Yukawa coupling ratio, *not* the VEV ratio $\tan \beta = v_u/v_d$, which is a free parameter of the Higgs potential);
- the top–bottom hierarchy $m_t \gg m_b$ (up quarks couple at 2nd BCH order, down quarks at 3rd, due to the pure-symmetric structure of $[T_{B-L}, g]$);
- the Cabibbo chain $\sqrt{m_d/m_s} \approx \lambda_W = \tan \theta_0$ (1.3% match, three independent observables from one bracket projection).

The new direction: $\text{SU}(2)_R$ vacuum misalignment

The CKM mixing requires the up-type and down-type sectors to see *different* torsion condensates. In Pati–Salam, this arises when the $\text{SU}(2)_R$ breaking direction is misaligned with the $\text{SU}(4)$ breaking direction (T_{B-L}). The most general bi-doublet VEV is

$$\langle \Phi \rangle = \begin{pmatrix} v_u & 0 \\ 0 & v_d e^{i\alpha} \end{pmatrix},$$

where $v_u/v_d = \tan \beta$ is determined by the Higgs potential minimum (a free parameter of the scalar sector, not fixed by the bracket alone) and α is a relative CP-violating phase. The misalignment angle between the $\text{SU}(4)$ and $\text{SU}(2)_R$ breaking directions determines the

Cabibbo angle; the phase α determines the Jarlskog invariant. Computing this requires minimising the BCH-generated Higgs potential for the bi-doublet Φ and the right-handed triplet Δ_R — a well-defined variational problem whose coefficients are fixed by the Killing form.

C.1 Torsion coupling hierarchy and vacuum energy cancellation

On a basis adapted to the eigenvectors of $\text{ad}_{T_{B-L}}$, all 15 independent generators of $\mathfrak{sl}(4, \mathbb{R})$ fall into exactly *two* tiers under the torsion VEV $\|[T_{B-L}, X]\|^2$.

Two tiers, or 0:1:4? Both statements appear in this paper and they are consistent; they are statements about different generator sets, and the distinction is worth making explicit because it has caused confusion. The spectrum of $\text{ad}_{T_{B-L}}$ is $\{0^{(9)}, (\pm 4/3)^{(3)}\}$, so on *eigenvectors* the coupling $\|[T_{B-L}, X]\|^2$ takes exactly two values, 0 and $32/9$. That is the basis-independent statement, and it is the one tabulated below. The electroweak generators J_i, K_i of $\mathfrak{so}(4) = \mathfrak{su}(2)_L \oplus \mathfrak{su}(2)_R$, however, are *not* eigenvectors: each is an equal-weight combination of one Tier-0 and one Tier-2 generator, for instance $J_1 = (A_{01} + A_{23})/2$ with $\|[T_{B-L}, A_{01}]\|^2 = 0$ and $\|[T_{B-L}, A_{23}]\|^2 = 32/9$. Each therefore inherits

$$\|[T_{B-L}, J_i]\|^2 = \frac{1}{4} \cdot \frac{32}{9} = \frac{8}{9},$$

and the three physically-named sectors — unbroken, electroweak, colour-lepton — carry couplings $0 : 8/9 : 32/9 = 0 : 1 : 4$. The 0:1:4 ratio is thus a real statement about the electroweak-adapted generators, not a third eigenvalue; it is reported as such in the Predictions section, where it governs the GW-matter coupling hierarchy. Both readings are verified in exact rational arithmetic by `scripts/generate_paper_figures.py` (function `torsion_couplings`) and by `code/acs_codebase/extras/torsion_hierarchy.py`.

Tier 0 (zero coupling, 9 generators): the three Cartan generators H_1, H_2, H_3 ; the within-colour rotations A_{01}, A_{02}, A_{12} ; and the within-colour boosts S_{01}, S_{02}, S_{12} . These are the unbroken generators after $\text{SU}(4) \rightarrow \text{SU}(3) \times \text{U}(1)_{B-L}$.

Tier 2 (coupling $32/9$, 6 generators): the colour-lepton antisymmetric generators A_{03}, A_{13}, A_{23} and their symmetric partners S_{03}, S_{13}, S_{23} . These connect colour indices to the lepton index, with eigenvalue gap $1/3 - (-1) = 4/3$. See Figure 8.

Vacuum energy cancellation (theorem). The torsion-weighted vacuum energy density is

$$\rho_{\text{vac}}^{(\text{bosonic})} = \sum_{X \in \mathfrak{sl}(4)} \|[T_{B-L}, X]\|^2 \cdot K(X, X),$$

where $K(X, X) = 8 \text{Tr}(X^2)$ is the Killing form. For each Tier-2 pair (A_{i3}, S_{i3}) :

$$\|[T, A_{i3}]\|^2 = \|[T, S_{i3}]\|^2 = \frac{32}{9}, \quad K(A_{i3}) = -16, \quad K(S_{i3}) = +16.$$

Product per pair: $(-512/9, +512/9)$. Net: $+512/3 - 512/3 = 0$ **exactly**, verified in exact rational arithmetic (SymPy, zero floating-point involved). The 9 Tier-0 generators contribute nothing (zero coupling).

This is *not* supersymmetry (which cancels bosons against fermions). It is the Palatini Form–Function pairing: for every connection-sector generator A_{i3} with negative Killing form, there exists a vierbein-sector partner S_{i3} with identical torsion coupling and positive Killing form. The cancellation is topological — it depends on the fiber structure, not on the base-space geometry, and is therefore protected under gravitational wave perturbations. (Verified: `torsion_higgs_vacuum.py`, `acs_vacuum.py`.)




Figure 8: The torsion coupling hierarchy of $\mathfrak{sl}(4, \mathbb{R})$. Left: on the eigenvector-adapted A/S basis all 15 generators fall into two tiers under $\|T_{B-L}, X\|^2$: Tier 0 (zero coupling, 9 generators including the colour $\mathfrak{su}(3)$ block) and Tier 2 (coupling $32/9$, 6 generators connecting the colour and lepton sectors). Right: the same couplings read on the electroweak-adapted generators, where J_i, K_i are equal-weight mixtures of a Tier-0 and a Tier-2 generator and so carry $\frac{1}{4} \cdot \frac{32}{9} = \frac{8}{9}$, giving the ratio 0 : 1 : 4. The two readings are consistent: $8/9$ is not a third eigenvalue. The Palatini pairing ensures that each antisymmetric generator A_{i3} with $K = -16$ is matched by a symmetric partner S_{i3} with $K = +16$, yielding exact cancellation of the torsion-weighted vacuum energy.

Scope of the cancellation. The cancellation operates at the level of the torsion-weighted Killing form — the algebraic contribution to the vacuum energy from the $\mathfrak{sl}(4)$ fiber. It does *not* cancel the Coleman–Weinberg effective potential (loop corrections), the Higgs potential at its minimum ($V_{\min} \sim -10^8 \text{ GeV}^4$), or the fermionic vacuum energy ($\sim +10^7 \text{ GeV}^4$). The cancellation removes the Planck-scale bosonic contribution ($\sim 10^{74} \text{ GeV}^4 \rightarrow 0$), reducing the cosmological constant problem from 10^{121} to $\sim 10^{55}$ orders of magnitude [14].

Photon and graviton protection. The photon (T_{B-L} direction) has zero torsion coupling at all BCH orders: $[T_{B-L}, T_{B-L}] = 0$, $[T_{B-L}, \text{SU}(3)] = 0$, and the 3rd-order bracket $[[T_{B-L}, g], T_{B-L}]$ is pure antisymmetric while T_{B-L} is symmetric, giving zero inner product. Gravitational waves propagate at c exactly: the contortion $K_{\nu\mu\rho}$ is antisymmetric in (ν, μ) , so $K_{\nu\mu\rho} k^\nu k^\mu = 0$ for any wavevector k . Both GW polarisations ($h_+, h_\times$) are symmetric rank-2, so no birefringence arises. (Verified: `torsion_causal.py`, `gw_torsion.py`.)

C.2 The irreducible boundary

The Pati–Salam Higgs potential $(\Phi + \Delta_R)$ has approximately 10 independent quartic couplings. The bracket algebra and consistency filters described in §6.5.1 together fix or exclude seven of them, leaving four free parameters in the minimal bi-doublet (Branch A).

Locked by Palatini bracket structure (5): $\lambda_\Phi = 2\sqrt{3}/27$ (Koide projection), the Palatini pairing $2\rho_1 + \rho_2 = 16/9$, the gauge coupling locking $g_4 = g_L = g_R = 4/3$, the Yukawa ratio $\tilde{h}/h = 2/3$, and the generation count $N_{\text{gen}} = 3$ from Jacobi truncation of the BCH at order three.

Forbidden by representation theory (1): $\alpha_2 = 0$ identically (Phase 50, §6.5.1): the cross-coupling $\text{Tr}(\Phi^\dagger T^A \Phi) \text{Tr}(\Delta_R^\dagger T^A \Delta_R)$ vanishes because $\Phi \sim (1, 2, 2)$ is an $\text{SU}(4)$ singlet.

Excluded at tree level by flavor phenomenology (1): $\beta_c = 0$ at tree level (Phases 51 and 52, §6.5.1): a non-zero β_c forces $\tan \beta = \pm 1$, which by the no-go identity (36) gives $|V_{\text{CKM}}| = \#$, contradicting observed quark mixing.

Free at tree level (3): ρ_1 (the Δ_R self-coupling, $(0, 8/9)$), α_1 (the diagonal Φ – Δ_R cross-coupling, $|\alpha_1| \lesssim 0.81$), and $\tan \beta$ (the bi-doublet VEV ratio; its tree-level value is undetermined due to Phase 51’s $\beta_c = 0$ flat direction. This flat direction is perturbatively protected by the Pati–Salam gauge structure: the ratio $\tilde{h}/h = 2/3$ is an RG invariant at one loop (both Yukawa couplings share identical gauge representations and therefore identical anomalous dimensions), two-loop QCD corrections are multiplicative and preserve the monotonicity of $V(\beta)$, and threshold corrections at v_R are β -independent at leading order in κ/v_R . Consequently, $\tan \beta$ cannot be fixed within the minimal bi-doublet sector at perturbative level; resolving it requires either a non-minimal Higgs sector or non-perturbative effects).

Calibrations (2): v and v_R .

Total inputs. With 2 calibrations and 4 free parameters, the Branch A model has **6 total inputs** versus the SM’s 19+, a reduction factor of approximately $3.2\times$. The Branch B alternative (extending the Higgs sector with $\Sigma \sim (15, 1, 1)$ to keep β_c at tree level) costs one additional parameter, raising the count to 7. Branch A is therefore the minimal viable embedding.

Falsifiability. The four free parameters are accessible to experiment: $\tan \beta \rightarrow m_t/m_b$ (perturbatively protected in the minimal sector; a non-minimal Higgs representation such as $\Sigma \sim (15, 1, 1)$ would introduce flavour-dependent threshold corrections capable of lifting the flat direction), $\alpha_1 \rightarrow$ Higgs-triplet mixing (LHC searches), $\rho_1 \rightarrow$ extra-Higgs spectrum (collider phenomenology), $v_R \rightarrow$ proton-decay rate.

Arithmetic lattice structure. We record a structural observation about the ratios collected above. The native dimensionless quantities of the bracket construction — the T_{B-L} charge ratio $1/3$, the adjoint eigenvalue $4/3$, the saturation coefficient $\text{ad}_{T_{B-L}}^3 = (16/9)\text{ad}_{T_{B-L}}$, the Yukawa ratio $\tilde{h}/h = 2/3$, and the multiplicities 3 and 9 — all lie in the multiplicative subgroup $\langle 2, 3 \rangle \subset \mathbb{Q}^\times$. The mechanism is mechanical: $N_{\text{color}} = 3$ supplies the prime 3, the quark–lepton charge difference $|q_q - q_\ell| = 4/3$ supplies the $2^2 \cdot 3^{-1}$ structure, and squaring gives $16/9 = 2^4 \cdot 3^{-2}$. No prime greater than 3 enters the construction.

This lattice is the same one that underlies Pythagorean tuning, where musical intervals are exactly the rationals in $\langle 2, 3 \rangle$ [45]. The mechanistic origins differ — Pythagorean tuning generates the lattice from octave equivalence and the third harmonic — but the resulting arithmetic structure coincides. The property is specific to the minimal Pati–Salam embedding: the Georgi–Glashow $\text{SU}(5)$ construction introduces the prime 5 through the $3+2$ trace split (eigenvalue gap $5/3$), and $\text{SO}(10)$ inherits this through its rank-5 Cartan structure. The companion note [46] contains the explicit comparison and an empirical test of whether the broken-phase mass spectrum and CKM mixing carry any imprint of the lattice; the test returns a clean negative result, placing the observation firmly as an ultraviolet algebraic artifact. We record it here as a structural feature of the algebraic sector with no phenomenological consequence claimed.

Grading selection from adjoint spectral structure. A related structural observation concerns the $\mathbb{Z}/2$ grading of the internal carrier space $\mathbb{R}^4$ of $\mathfrak{sl}(4, \mathbb{R})$. In the companion note [47], we show that the functional $\tilde{\mathcal{S}}_g[P] = \text{Tr}(\sigma_P \cdot g(\text{ad}_T))$ over symmetric involutions P , for any even $g \geq 0$ with $g(0) = 0$ and $g(\lambda) > 0$, is uniquely minimised when the P -negative eigenspace coincides with the minority eigenvalue cluster of T_{B-L} . For T_{B-L} this selects the $(3, 1)$ grading $P = \text{diag}(1, 1, 1, -1)$. The mechanism: the six quark–lepton mixing generators (the $\pm 4/3$ eigenspace of $\text{ad}_{T_{B-L}}$) constitute the algebra’s dynamically active sector, and the minimum forces all of them into the parity-odd subspace — which requires exactly the 3+1 split. The proof reduces to a max-cut problem on a complete bipartite graph (see [47], Theorem 1 and Lemma 1).

The result generalises to $\mathfrak{sl}(n, \mathbb{R})$ for $n = 3\text{--}6$ and across all four classical Lie algebra types. We emphasise, however, that the selected $(3, 1)$ grading acts on the *internal* $B-L$ charge space, not on physical spacetime. In light of the Coleman–Mandula theorem, identifying this internal decomposition with Lorentzian spacetime signature would require an additional bridging mechanism (soldering construction, pre-geometric reduction, or explicit evasion of the Coleman–Mandula hypotheses). See [47], §5 for a detailed discussion.

The ACS nesting chain: each solution becomes the next constraint




Figure 9: The ACS nesting chain (see Section D). Each layer is the solution that becomes the constraint for the next: gravity produces curvature via the Palatini bracket; curvature selects the colour algebra $\mathfrak{su}(3)$ as the unique closure attractor; the Koide projection and Higgs potential generate the mass spectrum; three generations follow from Jacobi truncation of the BCH; and falsifiable predictions emerge at each level. Within Branch A (minimal bi-doublet) the framework requires 2 calibrations (v, v_R), forbids 1 cross-coupling (α_2 , representation theory), excludes 1 at tree level (β_c , flavor no-go), and leaves 3 parameters free ($\rho_1, \alpha_1, \tan \beta$ radiative), for a total of 6 inputs (vs. 19+ in the SM).

D Geometric Vision (Plain-Language Summary)

Start with gravity alone: the shape of spacetime, encoded in a vierbein (the ruler) and a spin connection (the compass). These two fields are codependent and asymmetric. The asymmetry generates structure: the BCH expansion produces three layers that correspond to the Higgs mechanism (1st order), gauge bosons (2nd order, $\mathfrak{su}(3)$ closure attractor), and three generations of fermions (3rd order, Jacobi truncation). The full nesting runs: vibrations $\rightarrow$ momentum $\rightarrow$ colour $\rightarrow$ confinement $\rightarrow$ mass $\rightarrow$ gravity.

E Experimental Tests

The framework makes four falsifiable predictions, ranked by near-term testability.

1. 49 keV sterile neutrino (highest priority). The see-saw product formula (23) predicts $M_R \approx 49$ keV, producing an X-ray decay line at $E_\gamma = 24.5$ keV with mixing $\theta^2 \sim 10^{-6}$. XRISM (launched 2023) can reach flux sensitivity $\sim 10^{-7}$ ph/cm²/s in a 1 Msec exposure on Perseus. Detection at 24.5 keV would validate the geometric suppression chain; detection at a different energy would falsify $M_R = 49$ keV; non-detection constrains the production mechanism but not the mass. Timeline: 2–5 years.

2. $\theta_{\text{QCD}} = 0$ without axion. The real $\mathfrak{sl}(4)$ structure gives $\theta = 0$ as a theorem. The n2EDM experiment at PSI (target $d_n \sim 10^{-27}$ e·cm, ~ 2027) improves the current limit by 10 $\times$. If d_n is found above 10^{-27} , the theorem is under pressure. If no axion is found in the full PQ mass range (1 μ eV to 1 meV) by ADMX/CASPER, the ACS alternative to the strong CP problem gains support.

3. Large $\tan \beta$ ($\gg 1$). The Yukawa ratio $\tilde{h}/h = 2/3$ requires $\tan \beta \approx 30$ –60 for $m_t/m_b \sim 40$. LHC Run 3 and HL-LHC probe heavy Higgs bosons with enhanced $b\bar{b}/\tau\tau$ couplings up to $m_H \sim 2$ TeV. Current $B \rightarrow \tau\nu$ data show a 1.7 σ excess consistent with $\tan \beta \sim 30$ –50 for $m_{H^\pm} \sim 500$ GeV.

4. GW-spin coupling hierarchy 0:1:4 (long-term). Across the unbroken, electroweak and colour-lepton generators (0 : 8/9 : 32/9; see the torsion-hierarchy section). The torsion correction to GW-matter coupling is suppressed by $(v_R/M_{\text{Pl}})^2 \lesssim 10^{-6}$, below the sensitivity of current and near-future detectors. This prediction requires 3rd-generation instruments (Einstein Telescope, Cosmic Explorer, ~ 2035 +) and is currently untestable in practice.

Amendment: scope and evidence, 24 September 2026

The algebraic closure calculations do not establish a unique microscopic dynamics. For the stated smooth, metric-compatible Palatini action, writing the connection as its Levi-Civita part plus contortion leaves the added torsion-squared term algebraic in contortion. It supplies no independent contortion wave operator. Also, a balanced vector in the fundamental color triplet is not a singlet: vanishing Cartan expectations do not imply that the generators annihilate the state. These two inferences are withdrawn. The complete scalar action and physical vacuum still require coefficients, kinetic normalization and a selection rule.

For the complete-action and microscopic-carrier audit, see the companion papers *Complete Scalar Actions and Identifiability Boundaries in the ACS Framework* and *Condensate Binding, Charge Protection, and the Boundary of Particle Identification* in this repository's `papers/updates/` directory. The dated publication record separates fresh portable checks from preserved archival results; it does not claim a complete rerun of every historical experiment or a proof of RH.

References

- [1] T. Schreiber, *Measuring information transfer*, Physical Review Letters **85**(2), 461–464 (2000).
- [2] T. Bossomaier, L. Barnett, M. Harré, and J. T. Lizier, *An Introduction to Transfer Entropy: Information Flow in Complex Systems*, Springer, Cham (2016).
- [3] A. Kraskov, H. Stögbauer, and P. Grassberger, *Estimating mutual information*, Physical Review E **69**(6), 066138 (2004).
- [4] E. B. Dynkin, *Calculation of the coefficients in the Campbell–Hausdorff formula*, Doklady Akademii Nauk SSSR **57**, 323–326 (1947).
- [5] E. B. Dynkin, *Maximal subgroups of the classical groups*, Trudy Moskovskogo Matematicheskogo Obshchestva **1**, 39–166 (1952); English translation in AMS Translations, Series 2, **6**, 245–378 (1957).
- [6] B. C. Hall, *Lie Groups, Lie Algebras, and Representations: An Elementary Introduction*, 2nd ed., Springer Graduate Texts in Mathematics **222** (2015).
- [7] L. Onsager and S. Machlup, *Fluctuations and irreversible processes*, Physical Review **91**(6), 1505–1512 (1953).
- [8] M. I. Freidlin and A. D. Wentzell, *Random Perturbations of Dynamical Systems*, Springer, New York (1984).
- [9] Ya. G. Sinai, *Gibbs measures in ergodic theory*, Russian Mathematical Surveys **27**(4), 21–69 (1972).
- [10] D. Ruelle, *A measure associated with axiom-A attractors*, American Journal of Mathematics **98**(3), 619–654 (1976).
- [11] R. Bowen, *Equilibrium States and the Ergodic Theory of Anosov Diffeomorphisms*, Lecture Notes in Mathematics **470**, Springer (1975).
- [12] T. W. B. Kibble, *Lorentz invariance and the gravitational field*, Journal of Mathematical Physics **2**(2), 212–221 (1961).
- [13] D. W. Sciama, *The physical structure of general relativity*, Reviews of Modern Physics **36**(1), 463–469 (1964).
- [14] F. W. Hehl, P. von der Heyde, G. D. Kerlick, and J. M. Nester, *General relativity with spin and torsion: Foundations and prospects*, Reviews of Modern Physics **48**(3), 393–416 (1976).
- [15] B. A. Bilby, R. Bullough, and E. Smith, *Continuous distributions of dislocations: A new application of the methods of non-Riemannian geometry*, Proceedings of the Royal Society A **231**(1185), 263–273 (1955).
- [16] A. Palatini, *Deduzione invariantiva delle equazioni gravitazionali dal principio di Hamilton*, Rendiconti del Circolo Matematico di Palermo **43**, 203–212 (1919).

- [17] H. L. Montgomery, *The pair correlation of zeros of the zeta function*, in *Analytic Number Theory*, Proc. Sympos. Pure Math. **24**, AMS, Providence, RI, pp. 181–193 (1973).
- [18] A. M. Odlyzko, *On the distribution of spacings between zeros of the zeta function*, Mathematics of Computation **48**(177), 273–308 (1987).
- [19] H. Davenport, *Multiplicative Number Theory*, 3rd ed., Springer Graduate Texts in Mathematics **74** (2000).
- [20] M. F. Atiyah and I. M. Singer, *The index of elliptic operators on compact manifolds*, Bulletin of the American Mathematical Society **69**(3), 422–433 (1963).
- [21] A. B. Zamolodchikov, “Irreversibility” of the flux of the renormalization group in a 2D field theory, JETP Letters **43**(12), 730–732 (1986).
- [22] Z. Komargodski and A. Schwimmer, *On renormalization group flows in four dimensions*, Journal of High Energy Physics **2011**(12), 099 (2011).
- [23] J. C. Pati and A. Salam, *Lepton number as the fourth “color”*, Physical Review D **10**(1), 275–289 (1974).
- [24] A. Ashtekar, *New variables for classical and quantum gravity*, Physical Review Letters **57**(18), 2244–2247 (1986).
- [25] J. F. Barbero G., *Real Ashtekar variables for Lorentzian signature space-times*, Physical Review D **51**(10), 5507–5510 (1995).
- [26] G. Immirzi, *Real and complex connections for canonical gravity*, Classical and Quantum Gravity **14**(10), L177–L181 (1997).
- [27] M. Domagala and J. Lewandowski, *Black-hole entropy from quantum geometry*, Classical and Quantum Gravity **21**(22), 5233–5243 (2004).
- [28] K. A. Meissner, *Black-hole entropy in loop quantum gravity*, Classical and Quantum Gravity **21**(22), 5245–5251 (2004).
- [29] C. Rovelli, *Quantum Gravity*, Cambridge University Press (2004).
- [30] T. Thiemann, *Modern Canonical Quantum General Relativity*, Cambridge University Press (2007).
- [31] J. Engle, E. Livine, R. Pereira, and C. Rovelli, *LQG vertex with finite Immirzi parameter*, Nuclear Physics B **799**(1–2), 136–149 (2008).
- [32] R. Connelly and W. Whiteley, *Second-order rigidity and prestress stability for tensegrity frameworks*, SIAM Journal on Discrete Mathematics **9**(3), 453–491 (1996).
- [33] C. W. J. Granger, *Investigating causal relations by econometric models and cross-spectral methods*, Econometrica **37**(3), 424–438 (1969).
- [34] L. Barnett, A. B. Barrett, and A. K. Seth, *Granger causality and transfer entropy are equivalent for Gaussian variables*, Physical Review Letters **103**(23), 238701 (2009).
- [35] B. Wallace, *The Inversion Arc: Holographic Resolution in Asymmetric Coupled Systems*, 2026.

- [36] S. F. King, *A model of quark and lepton mixing*, JHEP **01**, 119 (2014), arXiv:1311.3295.
- [37] S. Coleman and J. Mandula, *All possible symmetries of the S matrix*, Physical Review **159**(5), 1251–1256 (1967).
- [38] S. Alexander, B. Alexandre, M. Fine, J. Magueijo, and E. Nākums, *GraviGUT unification with revisited Pati–Salam model*, arXiv:2510.11674 [hep-th] (2025).
- [39] A. H. Chamseddine, A. Connes, and W. D. van Suijlekom, *Grand unification in the spectral Pati–Salam model*, JHEP **11**, 011 (2015), arXiv:1507.08161.
- [40] F. Nesti and R. Percacci, *Graviweak unification*, Journal of Physics A **41**, 075405 (2008), arXiv:0706.3307.
- [41] M. Belfkir, M. A. Loualidi, and S. Nasri, *Fermion masses and mixing in Pati–Salam unification with S_3 modular symmetry*, Progress of Theoretical and Experimental Physics **2025**(3), 033B05 (2025).
- [42] M. Zantedeschi et al., *A minimal Pati–Salam theory: from cosmic defects to gravitational waves and colliders*, arXiv:2504.01893 [hep-ph] (2025).
- [43] B. Wallace, *Spectral Susceptibility and Renormalized Stability on the Riemann Critical Line*, 2026.
- [44] I. Lakatos, *Proofs and Refutations: The Logic of Mathematical Discovery*, Cambridge University Press, Cambridge (1976).
- [45] A. Barker, *Greek Musical Writings, Vol. II: Harmonic and Acoustic Theory*, Cambridge University Press, Cambridge (1989).
- [46] B. Wallace, *Pythagorean Arithmetic Structure in the Minimal Pati–Salam Bracket Algebra*, 2026.
- [47] B. Wallace, *Grading Selection from Adjoint Spectral Activity in $\mathfrak{sl}(n, \mathbb{R})$* , 2026.