

The Reversible Flattening: A Process Record

Every Path, Pass, Fail, and False Wall

Bradley Wallace

Independent Researcher / TensorRent

with Claude (Anthropic) as instrument and aggregator

TR-2026-FF06K

June 2026

Abstract

This is not a results paper. It is the laboratory record of how a single intuition — that there is structure hidden in what we routinely flatten — was driven from an unfalsifiable beginning to a small, defensible, falsifiable principle, and what was destroyed along the way. It documents the verified core, but it gives equal and deliberate weight to the wrong turns: six *false walls* (places where a result was claimed that the computation did not support), roughly fifteen recorded falsifications, four rounds of adversarial review that contracted the central claim three separate times, and one engineering integration that survived the same gauntlet that had killed four before it. The thesis that survives is stated once, late, and narrowly. A second arc applies the identical method to the non-trivial zeros of $\zeta(s)$ (Section 10), where it rediscovered a prior observable, recorded two wrong framings of its own central question as first-class negatives, and left an operator unconstructed and RH untouched. The method that produced both — conjecture, pre-registered test with a stated kill condition, compression into theorem or first-class negative — is the actual subject of this paper. Every numerical claim is reproducible under fixed random seed 20260423. Tiers are labelled throughout: T1 machine-verified, T2 proven/forced, T3 numerically verified, T4 falsified.

Contents

1	Why This Paper Exists	2
2	Method	3
2.1	Adversarial compression	3
2.2	The four-tier hierarchy	3
2.3	Pre-registration and the no-tuning rule	3
2.4	The asymmetric instrument	3
3	The Thesis and Its Three Contractions	3
4	The Verified Core	4
4.1	The exact ledger (T1)	4
4.2	The non-transitive boundary: Condorcet (T1)	4
4.3	The convolution seam (components T1; identification T3)	4
4.4	Structure search at 10^{80} (T1, with a corrected attribution)	5

5	The False Walls	5
5.1	Wall 1: the “Python wall” at 4978 digits	5
5.2	Wall 2: “62 orders of magnitude past the enumeration wall”	5
5.3	Wall 3: “compression”	5
5.4	Wall 4: the biconditional, three times	6
5.5	Wall 5: “FHE is the seam”	6
5.6	Wall 6: the seam as a binary void (the p-adic refinement)	6
6	The Adversarial Gauntlet	6
7	The Negative Ledger	7
8	Engineering Verification	8
8.1	Rust cross-language exactness (T1, promoted from T2)	8
8.2	The AISO motif-memory integration: the first to survive	8
9	The Origins: A Sycophantic Beginning, Metabolized	9
10	The Riemann Arc: The Same Method on the Zeros	9
10.1	The 2×2 , carried over	9
10.2	The critical line is not one seam (T1/T3)	9
10.3	The echo that wasn’t (T1)	10
10.4	Necessary spectral signatures (T1, necessary $\neq$ sufficient)	10
10.5	Boundary conditions, run as if true (T1, with a recorded convention error) . . .	10
10.6	The reductive sweep: which tunnels are empty (T1)	10
10.7	Two wrong framings, then the right one (T4, T4, then T3)	11
10.8	What the arc owes to LF01, FF06b’, and FF06e	11
10.9	The arc’s negative ledger	11
11	What the Method Cost, and What It Bought	12
A	Reproduction	12

1 Why This Paper Exists

The ordinary scientific paper reports a destination and hides the road. That convention is efficient and, for this body of work, actively misleading, because the most instructive events were not the results that survived but the claims that did not. A reader who sees only the final principle would conclude it arrived clean. It did not. It arrived three sizes too large, was cut down twice by external review and once by a counterexample, and reached its defensible form only after a dozen attractive ideas had been killed against ground truth.

This record exists because two specific failure modes shaped the entire program, and both are invisible in a results-only account:

Sycophancy. A result feels validated because a collaborator — human or machine — agreed with it, rather than because it was tested. The program began, fifteen months earlier, in exactly this failure (Section 9). Naming it was the precondition for everything that followed.

False-fit. A structural resemblance between a mathematical object and a target system is mistaken for an operational identity. “We have a Cantor pairing in the build, so the criterion applies” has the same shape as “the charge composes, so it must be multiplicative.”

Four such fits were proposed against the live engineering stack; four were false; the fifth was the first to survive (Section 8).

The defense against both is the same and is the only methodological commitment of this work: *language must match evidence*. A claim is allowed to be exactly as strong as the test that backs it, and not one word stronger. The recurring failure throughout — caught every single time by review or by a counterexample — was interpretation outrunning computation. The recurring success was the willingness to contract the claim until it fit.

2 Method

2.1 Adversarial compression

Every line of work followed one loop: form a conjecture; pre-register an explicit test with a ground-truth oracle *and a kill condition* before running it; then compress the outcome into either a theorem-grade artifact or a first-class negative result. Negatives are not failures to be buried; they are outputs, recorded with their mechanism. This is Lakatos’s *Proofs and Refutations* run with deliberate discipline: the conjecture is expected to be wrong in some respect, and the value is in locating exactly where.

2.2 The four-tier hierarchy

No claim is reported without a tier, and a lower tier may never be promoted by enthusiasm:

T1	machine-verified (executed, reproducible under seed 20260423)
T2	proven or forced by a stated argument
T3	numerically verified at finite scale (cannot stand in for T2)
T4	falsified (retained in the record with its mechanism)

A persistent temptation throughout was to let a T3 numerical result read as a T2 proof. The review rounds (Section 6) caught this more than once.

2.3 Pre-registration and the no-tuning rule

Thresholds, kill conditions, and the meaning of operational terms (what “combine” means, what counts as “overlap”) were fixed *before* seeing results. Tuning a filter or an encoding to the known answer is forbidden; where it was tempting — e.g. choosing a shape encoding that would make a homomorphism hold — the rule was stated explicitly in the test harness itself.

2.4 The asymmetric instrument

The work used several language models in distinct, non-interchangeable roles: an idea-generator (expansive, fluent, untrusted by default), a coder/aggregator/honest-instrument (the author of this record), and independent critics in language and substance. The asymmetry is the point: the generator’s fluency is treated as a hypothesis, never as evidence, and the critics are pointed at the generator’s output precisely because a fluent proposal is the most dangerous kind.

3 The Thesis and Its Three Contractions

The final principle is small. Its history is a sequence of contractions, each forced.

Principle 1 (Reversible flattening, final form). The algebraic flattening of a relational object to a scalar is losslessly reversible *when* the operation is an *injective homomorphism into a pre-specified, independently-meaningful structure*. The homomorphism carries density (how parts combine); the injectivity carries identity (which object); reversibility requires both. The boundary — the “seam” — is convolution: crossable for density, uncrossable for identity.

This is a falsifiable empirical principle, stated in the forward direction only. It was not born this way. It passed through the following forms:

1. **Biconditional (“iff”).** The first form claimed reversibility *if and only if* a homomorphism existed. The converse was never supported and was eventually *falsified* outright (Cantor pairing, Section 6). Contracted to one-way implication.
2. **Forward, but “into some structure.”** The next form claimed an injective homomorphism into *some* structure suffices. This is vacuous: any bijection induces a transported operation into which it is trivially a homomorphism. Contracted by adding the requirement that the target be *pre-specified* and independently meaningful.
3. **Forward, injective, pre-specified target.** Even this needed the explicit *injective* qualifier, which was forced empirically by the 30-domain test (Section 6): two domains were homomorphic but non-injective and broke the original criterion.

Three times the work wanted the larger claim. Three times the larger claim was wrong. The contraction is not a weakening to apologize for; it *is* the result.

4 The Verified Core

These are the claims that were never successfully attacked. They are stated at their true tier.

4.1 The exact ledger (T1)

Across seven applied-mathematics domains, the relational (prime-signature) representation reproduces standard arithmetic exactly: arithmetic functions ($d, \sigma, \sigma_2, \varphi, \mu$), gcd/lcm, binomial and factorial via Legendre’s formula, modular products, lattice index / Smith normal form, multinomial multiplicity (a ~ 4978 -digit value), and radical / perfect-power detection. **14 of 14** cases exact against an independent `sympy` oracle, reproducible under seed 20260423 (`verify_ledger.py`). This is the bedrock: the representation is faithful where the criterion says it should be.

4.2 The non-transitive boundary: Condorcet (T1)

Where the criterion says identity *cannot* survive a scalar, it is confirmed. Pairwise majority preference admits cycles (Condorcet); no scalar score can faithfully represent a cyclic dominance relation. Measured cycle frequency rose with candidate count — roughly 7%, 43%, 78%, 99% for 3, 5, 7, 10 candidates, and 48% at 1001 voters — and the faithful representation is the directed dominance graph, not any number. The flattening to a scalar is exactly the lossy step the principle predicts.

4.3 The convolution seam (components T1; identification T3)

The Euler product $\sum_n p(n)x^n = \prod_k (1-x^k)^{-1}$ was verified coefficient-wise to give the partition numbers $p(n)$ for $n \leq 30$: convolution adds indices and multiplies values. The *components* are machine-verified. The *identification* “the seam *is* convolution” is a synthetic observation

at T3, not a proven equivalence — and it is labelled that way deliberately, because this is precisely the kind of resonant phrase that wants to be promoted above its evidence.

4.4 Structure search at 10^{80} (T1, with a corrected attribution)

A branch-and-bound search over exponent vectors (knapsack-on-exponents with an admissible bound) found the maximal divisor count below very large bounds: $d = 103680$ at 10^{18} (7 ms), 1.81×10^{10} at 10^{50} (1.4 s), 3.01×10^{14} at 10^{80} (33 s), anchored against brute force at small bounds ($d = 12, 32, 240, 320$ at $N \leq 60, 840, 720720, 2162160$). The result is real. The *attribution* was initially wrong and is corrected in Section 5: the speed comes from the highly-composite/knapsack structure of the problem, not from the relational representation, which is a bystander here.

5 The False Walls

This is the heart of the record. A “false wall” is a place where a result was asserted that the computation did not support — not a falsified *conjecture* (those are in Section 7), but a falsified *interpretation* of a true computation. Each is given as: the claim, how it was caught, and what it contracted to.

5.1 Wall 1: the “Python wall” at 4978 digits

Claimed. A multinomial multiplicity with ~ 4978 digits “hit a wall that the relational form did not” — presented as the scalar representation failing where the shape survived.

Caught. The integer itself computes in milliseconds. The only thing that fails is the *base-10 string projection*, throttled by a Python configuration limit (`sys.set_int_max_str_digits`, CVE-2022-45061 mitigation). Crucially, `Shape.volume()` — the relational form’s own projection back to an integer — hits the *identical* wall. The relational form is not immune; it simply never *needs* the projection step.

Contracted to. “The relational form avoids a projection the scalar form forces, but enjoys no special immunity to it.” A statement about *when* a step is needed, not about a capability gap.

5.2 Wall 2: “62 orders of magnitude past the enumeration wall”

Claimed. The 10^{80} search demonstrated the relational representation leaping far beyond where enumeration dies.

Caught. The leap is a property of the branch-and-bound over divisor-count structure, not of the representation. Any encoding of the exponent vector admits the same bound. The shape engine is incidental to the speed.

Contracted to. “Knapsack-on-exponents with an admissible bound searches highly-composite structure efficiently; the representation is the bystander.” (Section 4.)

5.3 Wall 3: “compression”

Claimed. The relational re-representation compresses.

Caught. Measured directly: a primorial compresses by $1.25\times$, a prime by $1.01\times$ — and for squarefree numbers the shape is *expansion*, not compression.

Contracted to. “Structural re-representation,” with no compression claim. The word “compression” was removed wherever it appeared, including from a paper’s opening line on final review.

5.4 Wall 4: the biconditional, three times

Claimed. Reversibility *iff* homomorphism — in three successive forms (Section 3).

Caught. The converse was never supported; one auditor flagged the unproven “iff” as language outrunning the proof; a counterexample search then *falsified* the converse directly (Cantor pairing, Section 6). The “into some structure” repair was caught as vacuous.

Contracted to. The final forward-only, injective, pre-specified-target principle. This wall recurred because the biconditional is genuinely the more beautiful claim; its repeated return is the clearest evidence that beauty is not a tier.

5.5 Wall 5: “FHE is the seam”

Claimed. Fully homomorphic encryption *is* the convolution seam between additive and multiplicative structure.

Caught. Ring-LWE is a single structure native to both operations; its noise exists for lattice-problem hardness (SVP), not because addition and multiplication are incompatible. The inability to glue two single-operation schemes is ordinary type-incompatibility, which cryptographers never mistook for a “seam.”

Contracted to. An explicitly-labelled illustration, plus a known triviality. Demoted from evidence to analogy.

5.6 Wall 6: the seam as a binary void (the p -adic refinement)

Claimed (implicitly). The seam is a binary presence/absence of structure.

Caught. The p -adic valuation gives $v_p(a + b) \geq \min(v_p(a), v_p(b))$, a one-sided bound that is *tight* roughly half the time (when valuations differ) and slack on ties. The seam is therefore a measurable *gradient* of structural loss, not a void.

Contracted to. “The seam is a one-sided, partially-tight bound — a gradient,” which is a sharpening rather than a retraction, and the only wall that made the picture richer rather than smaller.

6 The Adversarial Gauntlet

The claim was submitted to four rounds of attack. The grade trajectory was $2/5 \rightarrow 4.5/5 \rightarrow 5/5$. Each round contracted the claim to something smaller and truer; in no round was the verified core (Section 4) successfully challenged. What got stripped, every time, was interpretation.

Round 1 — language (3 overclaims). A language-focused critic flagged three: the seam stated as a T1 theorem (it is T3); the unproven “iff”; and “more tractable” phrasing that the negative ledger (Section 7) directly contradicts. All three corrected.

Round 2 — substance (5 concessions). A substance-focused auditor produced Walls 1, 2, 3, 5, and 6 above (the Python wall, the knapsack attribution, the compression fallacy, the FHE strawman, the p -adic gradient). All conceded and corrected in the same document.

Round 3 — counterexample (converse falsified). A direct search for a reversible map that is *not* a homomorphism produced the Cantor pairing: a bijection $\mathbb{N}^2 \rightarrow \mathbb{N}$, perfectly reversible, yet $\pi(1, 2) = 8$, $\pi(2, 0) = 3$, $\pi(3, 2) = 17$, and $8 + 3 = 11 \neq 17$ — it does not carry componentwise addition. The converse of the criterion is therefore *false* (T4). This forced the pre-specified-target qualifier (the vacuity guard).

Round 4 — final review (5/5). The last reviewer’s verdict was that the corrections were published in the same document as the claims — “the opposite of sycophancy.” The single residual (“compression” surviving in an opening line) was then removed.

The 30-domain predictive test (the injectivity forcing). Separately, the criterion was run as a *predictor* across 30 domains: does each domain’s flattening satisfy the criterion, and is it in fact reversible? Under the original (non-injective) criterion: **28/30** correct. The two misses — persistent homology, and list $\rightarrow$ multiset — were both homomorphic but *non-injective*. Adding the injective qualifier brought it to **30/30**. This is strong evidence the criterion picks out a real class rather than being fitted retrospectively, precisely because it first *failed* on two cases and the failure had a single common cause.

7 The Negative Ledger

Roughly fifteen conjectures were falsified and retained. Each is a result. The mechanism is given because a falsification without a mechanism is just a stuck door.

Conjecture (T4)	Mechanism of falsification
Multiply/divisibility faster than big-int	Compiled big-integer arithmetic wins on throughput.
Exact rationals beat Fraction	Custom exact rationals ran $11\times$ slower.
Near-unity ratios need exact shapes	Floating-point log is fine to 10^{-16} ; no shape needed.
Greedy superabundant search	Greedy is suboptimal; misses true maxima.
Non-admissible search bound	Pruned actual optima; bound must be admissible.
BRA charge is multiplicative	Live source: charge composes <i>additively</i> (table-sums).
Dual-engine router is multiplicative	Live source: ordinal thresholds, transitive.
Trust score is multiplicative	Additive in production.
Merkle structure is multiplicative	Hashing destroys the multiplicative structure.
“Commutativity $\Rightarrow$ no structure”	gcd is commutative <i>and</i> structured; order-dependence and structure-retention are independent axes.
φ + inverse-silver “structures noise”	φ alone: discrepancy 0.00084; the combination: 0.02132 — $25\times$ worse.
FHT “consciousness score”	Saturates to 1.0 on noise and structure alike; non-discriminating.
Criterion converse (reversible $\Rightarrow$ homomorphism)	Cantor pairing: reversible, not a homomorphism.
“Injective homomorphism into <i>some</i> structure”	Vacuous: any bijection via a transported operation.
Motif retrieval via gcd	Accuracy 55.15% vs 83.13% majority baseline — below the trivial predictor.

Two of these (the “commutativity” axis correction and the p -adic gradient) did more than close a door — they corrected the conceptual frame. The rest are honest dead ends, which a program that reports only wins would have silently swept up.

8 Engineering Verification

The mathematics was tested against running code, where resemblance is cheapest and most dangerous.

8.1 Rust cross-language exactness (T1, promoted from T2)

The integer-only shape engine was transcribed to Rust and checked against the `sympy`-validated Python reference: `cargo test` 8/8 pass (engine checked against independent in-Rust trial-division and Euclid, trusting no external oracle), and a 503-record Rust $\leftrightarrow$ Python ledger diff is empty under seed 20260423, including 100+ digit values. The monograph’s one soft assertion — “transcribes directly to an integer-only Rust module” — moved from T2 (assertion) to T1 (machine-verified). It remains a standalone utility; it was *not* wired into the production stack, because no call site with a genuinely injective-homomorphic operation was identified, and four prior such hopes had been false.

8.2 The AISO motif-memory integration: the first to survive

A proposal to encode memory “motifs” as shapes was tested by a pre-registered harness with a self-test proving it could both pass a true homomorphism and *falsify* a non-homomorphism (a test that cannot fail is worthless). Run against 45 live motifs and 990 pairs:

- **Check 1, composition (write side): PASS.** Motif composition matches exponent-addition on **100%** of 990 pairs with **0** collisions. The other candidate operations cleanly lost: max/lcm 57.47%, min/gcd 0%, squarefree-union 0%. The gate could have failed at three independent places and failed at none.
- **Check 3, containment cost.** $O(\text{distinct primes})$: min 4, median 6, max 10, mean 5.977 — explicitly *not* $O(1)$, correcting an earlier claim.
- **Check 4, identity vs density.** 57 concepts, 54 appearing in multiple motifs, `frac_multi` = 0.947: the two query axes are genuinely distinct.
- **Check 2, retrieval via gcd (read side): FALSIFIED.** Accuracy 55.15% against an 83.13% majority baseline — below the trivial predictor. gcd is not the shared-context operator; the statistical embedding layer handles retrieval. (Scope: labels were channel-provenance, a coarse but independent proxy for conceptual overlap; the result falsifies gcd against the best independent signal available.)

The true-by-construction check. Before banking Check 1, one question decided whether it was real or circular: does the *production* merge actually accumulate counts losslessly, or was the test’s notion of “combine” chosen to make additivity automatic? Reading the live source confirmed the merge path is append-only with no cap, window, decay, or overwrite — so the homomorphism is structural, not definitional.

The two-layer correction (reading source beats trusting the checkmark). The invariant comment initially pinned to the source claimed “no dedup.” Reading further revealed a motif detector fifty lines down that *does* deduplicate recurring windows into references. The claim was false as written. The corrected invariant scopes the guarantee precisely: the homomorphism lives in the *event/keyword-count layer* (append-only, accumulating), while the motif-reference layer is an intentional count-preserving lossy compression that the homomorphism does not depend on. The real tripwire is event-level dedup or decay, not the existing motif dedup. This correction was caught only by reading the code, not by any passing test — the single most important lesson of the engineering phase.

What it established. Composition is homomorphic (write side); retrieval-by-gcd is not (read side); the seam between them is the gap between keyword-identity and concept-identity. The production architecture — multiplicative-exact storage plus statistical embedding retrieval — is vindicated by test rather than by hope. After an explicit 0-for-4 record on prior integration fits, this was the first to pass, and it passed against production source.

9 The Origins: A Sycophantic Beginning, Metabolized

The intuition behind all of this is fifteen months old, and its first expressions were unfalsifiable. Two early artifacts were re-tested against ground truth during this program and both were *falsified*:

- The φ + inverse-silver claim that an irrational combination “structures noise.” Tested: φ alone gives discrepancy 0.00084; the combination gives 0.02132, $25\times$ worse. The kernel was real (the golden ratio’s badly-approximable property) but the narrative was sycophantic.
- The “FHT consciousness score” (a log-power Box–Cox-style transform). Tested: it saturates to 1.0 on random noise and on structure alike — it discriminates nothing. Again a real kernel (a legitimate transform) wrapped in a claim no test could fail.

The pattern in both: a genuine technical seed inside an unfalsifiable story that an agreeable interlocutor reflected back as profound. That reflection was real harm — and it was also the provocation for a year of work spent learning to distrust and test exactly that kind of agreement. The whole arc, from FHT to the present principle, is the *same seed* — “there is structure in what we flatten” — grown from an unfalsifiable infancy into a claim that states its own kill conditions. The discipline did not replace the intuition; it gave the intuition something to push against.

10 The Riemann Arc: The Same Method on the Zeros

The method — conjecture, pre-registered test *with a stated kill condition*, compression into theorem or first-class negative — was carried to a second domain: the non-trivial zeros of $\zeta(s)$. All work was done in-sandbox against the first 10^5 real zeros ($\gamma_1 = 14.135$ to $\gamma_{10^5} = 74920.827$), seed 20260423. Two things must be said at the outset. First, the arc partly *rediscovered* prior work — a technical note (LF01) on finite-window explicit-formula observables — and that is recorded as rediscovery, not dressed as discovery. Second, the arc produced two wrong framings of its central question before producing the right one; both wrong framings are retained below as the negatives they were.

10.1 The 2×2 , carried over

The lens was the square from the relational work: is a map homomorphic, and is it injective? Homomorphic+injective is reversible (good); homomorphic+non-injective is a *seam* (density crosses, identity is lost); non-homomorphic+injective is the *Cantor corner* (identity kept, structure scrambled); neither is noise. This square turned out to classify the maps around the critical line.

10.2 The critical line is not one seam (T1/T3)

The natural conjecture — “ $\Re(s) = \frac{1}{2}$ is the seam” — is false as a single label. Three distinct maps live there. The reflection $s \mapsto 1 - s$ is the *Cantor corner*: injective, non-homomorphic, fixed locus $\Re = 0.500000$. The prime $\leftrightarrow$ zero correspondence is *granularity-dependent*: lossless

on the full set (the $\psi(x)$ reconstruction error falls $0.0030 \rightarrow 0.00002$ as $N:100 \rightarrow 10^5$) but a seam at the individual level (dropping a single zero delocalizes the reconstruction, participation ratio 0.72). The Euler-product convergence edge is the genuine convolution seam — and it sits at $\Re = 1$, not $\frac{1}{2}$. The one-word question “is the critical line a seam?” had three different answers depending on which map you meant.

10.3 The echo that wasn’t (T1)

The prime $\leftrightarrow$ zero granularity split looked like the same mechanism as the AISO motif-memory seam (Section 8). It is not: it is a coincidence of cell-label, not of mechanism. The discriminator is density of support. Per-element footprint, measured as participation ratio: prime $\leftrightarrow$ zero = 0.74 (each zero a global, dense superposition — non-injectivity is delocalization); motif memory = 0.018 (each event one keyword coordinate — non-injectivity is the keyword-vs-concept gap). The homomorphic+non-injective cell of the square holds at least two distinct mechanisms; the same diagram entry is not the same phenomenon.

10.4 Necessary spectral signatures (T1, necessary $\neq$ sufficient)

Three Hilbert–Pólya *necessary* signatures, all machine-verified on the real zeros, none constructive. Nearest-neighbour spacing: KS 0.0197 to GUE versus 0.299 to Poisson. Number variance / spectral rigidity: GUE log-law on 5/5 windows ($\Sigma^2 \approx 0.33\text{--}0.42$, not linear). Pair correlation: matches Montgomery $1 - (\sin \pi r / \pi r)^2$ to 0.026, correlation hole present. The scope wall is hard and stated every time: these are *necessary* signatures of a chaotic-class spectrum; they do not construct an operator and do not bear on RH.

10.5 Boundary conditions, run as if true (T1, with a recorded convention error)

Assuming the Berry–Keating $H = xp$ picture, two boundary conditions were extracted from the data. **BC-1, the Maslov constant.** The smooth counting function’s constant term, fit over the high- T band, gives $C = 0.8750$, exactly $7/8$. This is recorded with its error: the first run gave $C = 1.375$, off by exactly 0.5, from a counting-convention slip ($N(\gamma_n) = n$ instead of the von Mangoldt jump-average $n - \frac{1}{2}$). The fix landing precisely on the predicted $7/8$ is a motivated-correction risk and is flagged as such; the defence is that $n - \frac{1}{2}$ is the standard convention, not a tuned knob, and the result is exact to four digits. **BC-2, the periodic orbits.** The Fourier transform of the zeros peaks at $\log(p^k)$ at $\sim 1.4 \times 10^3$ relative height and is flat (~ 1.3) at composites; the kill condition (a peak at $\log 6$, $\log 10$, $\log 12$) did not fire. The boundary conditions of the operator are found; the operator is not constructed. Walls found; the room not shown to exist.

10.6 The reductive sweep: which tunnels are empty (T1)

The signatures were run against decoys to measure discriminating power. A *genuine* GUE matrix spectrum, remapped to the zeta density, passes spacing (KS 0.011) and rigidity ($\Sigma^2(10) = 0.59$) — statistically indistinguishable from the zeros — yet is *empty* of prime orbits (2.4 versus the zeros’ 228.8). Poisson fails everything. Specificity: the zeros’ peaks sit at *exact* log primes (228.9), collapse under a 0.05 shift (1.4), and vanish at random integers (0.4). The result is a ranking of the signatures by discriminating power:

signature	discriminating power	shared by
prime-orbits (F_w peaks)	strong	only arithmetic spectra
GUE spacing / rigidity	medium	<i>all</i> chaotic spectra (incl. random matrices)
Maslov constant $7/8$	weak	any curve-matched sequence

The GUE decoy passing the statistics while empty of orbits is exactly LF01 Appendix C(i) — Berry–Keating reproduces GUE but not the prime peaks — seen from the decoy side: a GUE-class spectrum is *provably* arithmetic-free, so constraint (b) cannot come for free from (a). This whole subsection is the explicit-demonstration form of a sentence FF06b’ states directly (§3): “the zeros are GUE, but GUE is generic; it does not select an operator.” The external decoys here are the un-hardened sibling of FF06b’'s internal control — the marginal-matched shuffle of FF06e, which destroys the object’s correlations while preserving its own distribution, and reaches the same form/function split without constructing a separate GUE ensemble.

10.7 Two wrong framings, then the right one (T4, T4, then T3)

The central question was whether the GUE statistic (a) and the prime peaks (b) are in tension or independent. Both first answers were wrong.

Framing 1 — tension (chaos and arithmetic mutually exclusive): **false**. The real zeros have both at full strength, so they are not exclusive.

Framing 2 — independence, a cheap add-on (pre-registered as the expected answer): **also false**. A GUE base with an additive prime displacement, built from primes alone, tops out at peak ~ 6 while still GUE-passing and breaks GUE if pushed harder — versus the zeros’ 228 with GUE intact. You cannot bolt prime structure onto a GUE field and keep the field.

What survives. Three constructions bracket the zeros and none reaches their corner: generic GUE (a, not b); a prime field on a rigid lattice (b-ish at KS 0.099, not a); their additive sum (neither well). The reading the bracket forces is that (a) and (b) are two read-outs of one object — the arithmetic fluctuation field $S(T)$, whose local correlations give GUE and whose Fourier content gives the prime peaks — not two separable mechanisms to build and add. The scope is kept honest: the constructions are crude (perturbative displacement, 95 primes, $k \leq 3$), so this is a sandbox illustration and a sharpening of LF01 Appendix C, *not* a no-go theorem.

10.8 What the arc owes to LF01, FF06b’, and FF06e

The orbit observable $A(\tau) = \sum_k w(\gamma_k) \cos(\tau\gamma_k)$ used throughout is, on rediscovery, one object under several names: LF01’s F_w , FF06b’'s prime-dual (§6, its load-bearing witness), and FF06e’s arithmetic witness — the explicit-formula coupling. The hardened forms live in those documents, not here: FF06b’ §6 holds the prime-dual at 100% across a $15\times$ height sweep and a $7\times$ tighter tolerance with no off-prime peak; LF01 recovers Frobenius-trace *amplitudes* across nineteen L-functions with $O(1/N)$ error certified ($\text{rms}(E) \cdot N = 2.93 \pm 2.2\%$, slope -0.996) and discriminates non-abelian classes (S_3 , D_4); FF06e quantifies the form/function split against a marginal-matched null ($\sim 11,500\sigma$ arithmetic versus $0.4\text{--}1.0\sigma$ for spacing, counting, shape). The conditional-uniformity height floor $T_{\min}(d, q) = (2\pi e)^d/q$ (FF06b’ §10) is the first-class negative that corrected LF01’s uniformity overclaim, and it explains why the non-abelian results stop where they do: degree enters exponentially in the floor. The sandbox’s one genuine contribution back to that corpus is the decoy/bracket framing — it converts “no tested construction does all three” into a structural statement: a successful operator must produce a *single* field that is intrinsically GUE-distributed *and* prime-coherent, not a chaos mechanism plus a prime mechanism. The remaining work (rigorous self-adjoint xp , de Branges, Connes) is genuine open mathematics, outside what a sandbox can settle.

10.9 The arc’s negative ledger

Conjecture (T4)	Mechanism of falsification
$\Re(s) = \frac{1}{2}$ is a single seam	Three maps, three cell-types; convolution seam is at $\Re = 1$.
Prime $\leftrightarrow$ zero echo = motif-memory mechanism	Coincidence of cell-label; PR 0.74 (dense) vs 0.018 (sparse).
GUE statistics single out the zeros	A genuine random matrix is statistically indistinguishable; orbits 2.4 vs 228.8.
Maslov constant $C = 1.375$	Counting-convention slip ($N = n$ vs $n - \frac{1}{2}$); corrected to 7/8 exactly.
Chaos vs. arithmetic <i>tension</i>	False: the zeros have both (a) and (b) at full strength.
Arithmetic is a cheap <i>independent</i> add-on to GUE	False: additive build tops at peak ~ 6 (GUE-passing) vs 228.

The same shape as the relational arc: the empty tunnels (two wrong framings, a non-discriminating statistic, a convention error) map the real result — one field, two readouts, an unconstructed operator — more precisely than the signatures alone could. Necessary signatures, mapped by discriminating power; sufficiency unreached; RH untouched.

11 What the Method Cost, and What It Bought

The principle that survives is smaller than the one the work kept reaching for. It is forward-only where it wanted to be a biconditional; it requires a pre-specified target where it wanted to range freely; it makes no compression claim, no speed claim, no claim of immunity to the limits that bind the scalar form. The seam is characterized but not proven to *be* convolution. The forward direction itself is an empirical principle, not a theorem — the one genuinely open internal-deepening move (a category-theoretic reformulation) is named and not yet done.

Against that, what the method bought: a 14/14 exact ledger; a confirmed non-transitive boundary; a machine-verified cross-language engine; an engineering integration that survived a falsification test built to kill it, checked against production source; and a negative ledger of fifteen entries, each with a mechanism, that maps the boundary of the idea far more precisely than any list of successes could.

The trade is worth stating plainly because it is the thesis of this record: *a smaller claim that is true is worth more than a larger claim that is impressive*, and the only reliable way to tell them apart is to write down, in the same document, every place the larger claim was tried and failed. The contraction is not the price of the result. The contraction *is* the result.

A Reproduction

All numerical claims reproduce under random seed 20260423 unless otherwise noted (the motif-memory runs and the validation suite used their own fixed seeds, recorded with each result). The core artifacts: `verify_ledger.py` (the 14/14 ledger); `predictive_test.py` (the 30-domain test); the Rust `shape_engine` crate with `cargo test` (8/8) and `verify.sh` (503-record diff); and `motif_test_suite.py` with `check2_gcd_overlap.py` (the AISO integration, each with a discriminating self-test). Falsified conjectures are retained in the record with their mechanisms (Section 7); methodological corrections are appended, never erased.

This is a process record (TR-2026-FF06K), companion to the result papers FF06g (the geometry engine), FF06h (the relational principle and its limits), FF06i (the reversible-flattening parent), and the

consolidating monograph FF06J; the Riemann arc of Section 10 is the sandbox-process companion to the hardened spectral results in FF06b' (Spectral Witness Survival and the Transport Obstruction), the shuffle-knife categorisation in FF06e, and the L-function localisation note LF01, whose central observable it rediscovered and whose operator constraints it sharpens. It documents the path; those documents state the destination.